

Tutorial – PaaS

Case Study: Force.com application development platform

Week 6

School of Software

Faculty of Engineering and Information Technology

University of Technology Sydney



SCHOOL OF SOFTWARE

Short review of PaaS

- Overview of PaaS
 - Composed of **hardware resources** and **software resources** (such as operating systems and/or software frameworks) to ***build*** software applications;
 - Could be provisioned **on top of** IaaS;
 - **Shared hardware resources and software resource** usage among all PaaS consumers;
 - Major PaaS providers: **Google App Engine**, **Force.com platform**, and **Microsoft Azure**.

Case study for Force.com platform

- **Question 1:** What are multi-tenant applications?
- **Question 2:** What are the benefits of multi-tenant applications?
- **Question 3:** Architecture overview of Force.com
- **Question 4:** Data storage model of Force.com

Case study for Force.com platform

- **Question 1:** What are multi-tenant applications?
(refer to Section 2)
- **Question 2:** What are the benefits of multi-tenant applications?
(refer to Section 2 and 3)
- **Question 3:** Architecture overview of Force.com
(refer to Section 4 and 5)
- **Question 4:** Data storage model of Force.com
(refer to Section 6)

Q1: What are multitenant applications?

- Multitenant applications satisfy the needs of **multiple users or organizations** (in contrast to single tenant applications).
- In PaaS multi-tenant applications, all tenants **share the same underlying computing platform**
- Tenants operate in **virtual isolation** from one another, and their data and customizations remain **secure and insulated** from the activity of all other tenants.




Figure: Tenants sharing a single stack of computing resources

Q2: What are the benefits of multi-tenant applications?

- **Benefits to application providers:**

- Ease of carrying out application **update(s)**;
- Enabling and managing application delivery for multiple users using a **single stack of resources** (much easier and cheaper than maintaining a single stack of resources for each user);
- Provider **gather operational information** from the collective user population (which queries respond slowly, what errors happen more often, etc.)

- **Benefits to the end users:**

- Receive the applications at a (much) **lower price** than traditional mechanisms

Q3: Working of Force.com (at an abstract level)

- Core-application data and logic is stored as metadata
- End-user customizations (changes to existing core data, core logic, new data, and new logic), are stored as metadata (and not as views).
- Force.com provides clear separation between:
 - (a) Core application data and functionality (Force.com defined Data and Logic)
 - (a) User-defined application data and functionality (Custom Data and Logic)
- Force.com does not create an “actual” table in a database (in response to customizations), rather it’s runtime engine generates these views on the fly using the metadata

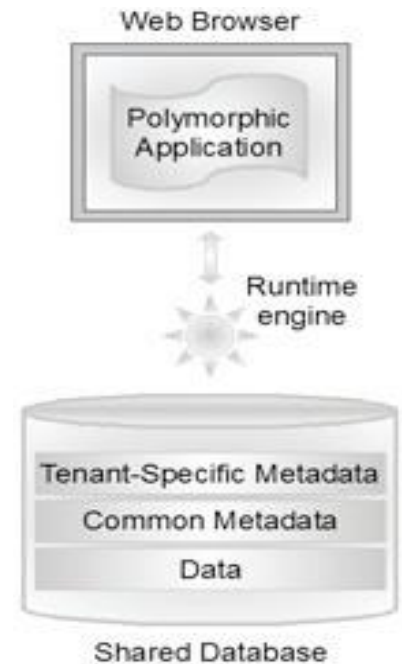


Figure: Force Multi-tenant Architecture

- Different users (tenants) see their customized (tenant-specific) view of the application

Q3: Architectural elements of Force.com

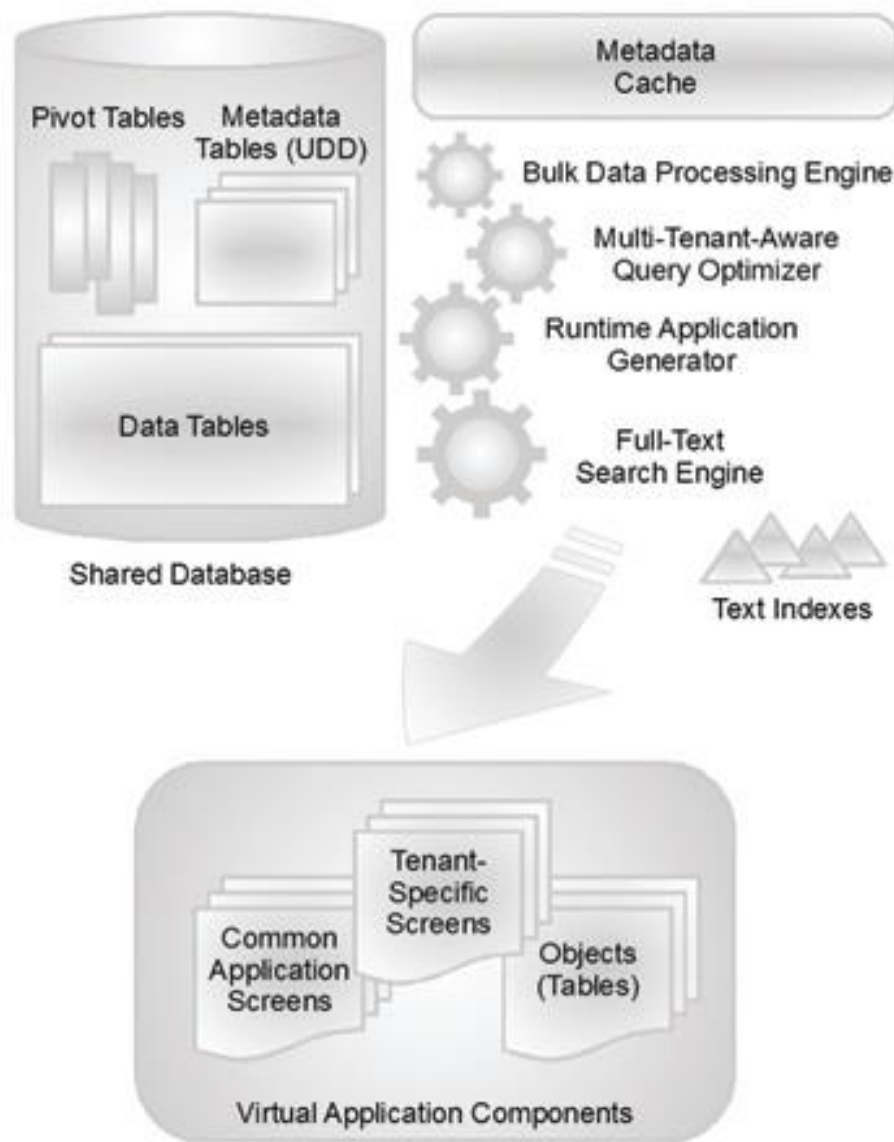


Figure: Architectural elements of Force.com platform

Q3: Architectural elements of Force.com

- Metadata cache:
 - Frequently accessed metadata (used to generate application components) is stored in metadata cache;
 - This improves performance and application response time.
- Bulk Data processing engine:
 - This engine carries out repetitive actions (such as bulk update, bulk retrieve, bulk delete etc...);
 - Has an intelligent in-built fault-recovery mechanism to re-execute the “update” or “save” operations after factoring out records that cause errors during the bulk update/save process

Q3: Architecture overview of Force.com

- Full text external search engine:
 - Force.com employs an **external search service**
 - This works independent of the other architectural elements
 - Tenant's data is **indexed** (and indexes are stored as a special set of pivot tables) for quick access and retrieval.
- **Runtime Application Generator:**
 - Runtime engine accesses the
 - core application metadata; (and)
 - custom metadata of the corresponding user.
 - Runtime engine **compiles them at run-time** to generate custom views for the corresponding user

Q3: Architecture overview of Force.com

- Multitenant-aware query optimizer
 - The underlying storage tables are partitioned based on Org ID
 - The query optimizer, considers which users is executing a given application function, and using intelligent algorithms determines which database partition needs to be queried

Q3: Architecture overview of Force.com

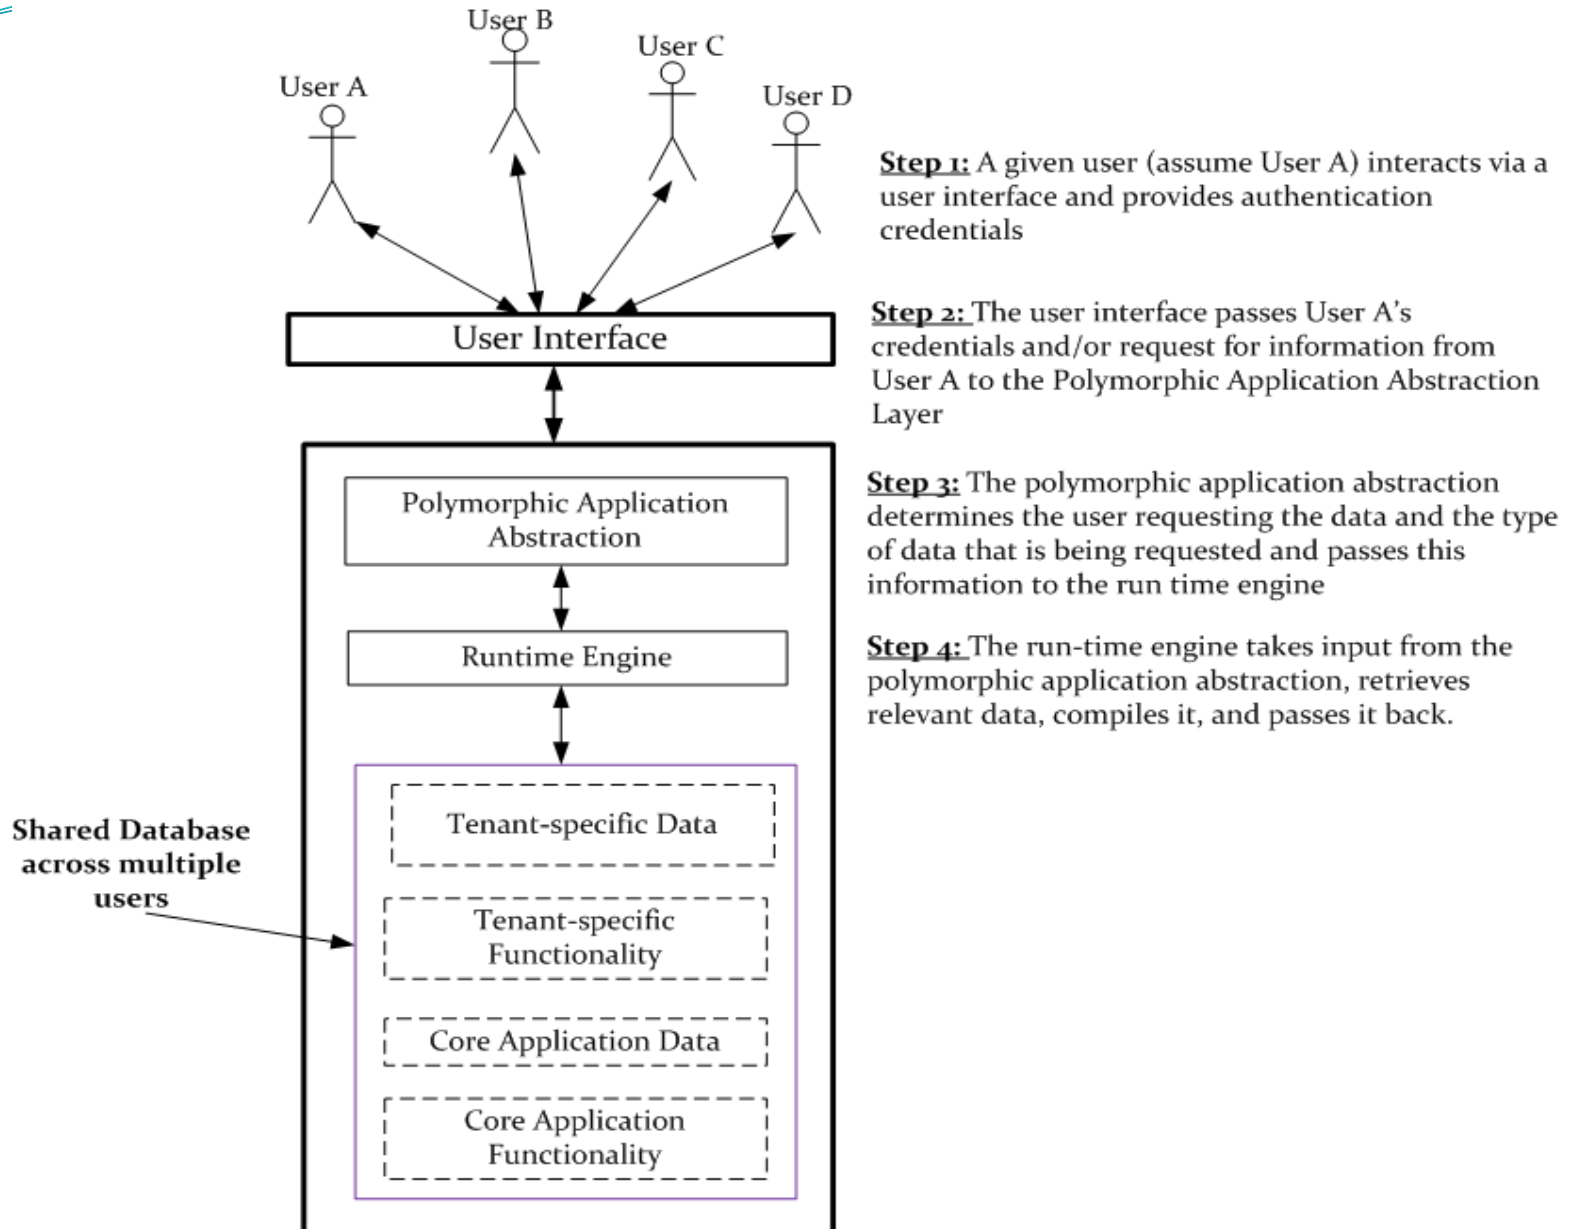


Figure: Step-wise working of Force.com

Q4: Data Storage Model of Force.com

- Force.com manages to deliver the views (including the customized user defined views) using a combination of:
 - Metadata tables
 - Data tables
 - Pivot Tables
- Joins between the above three tables are carried out at run-time to generate customized views.

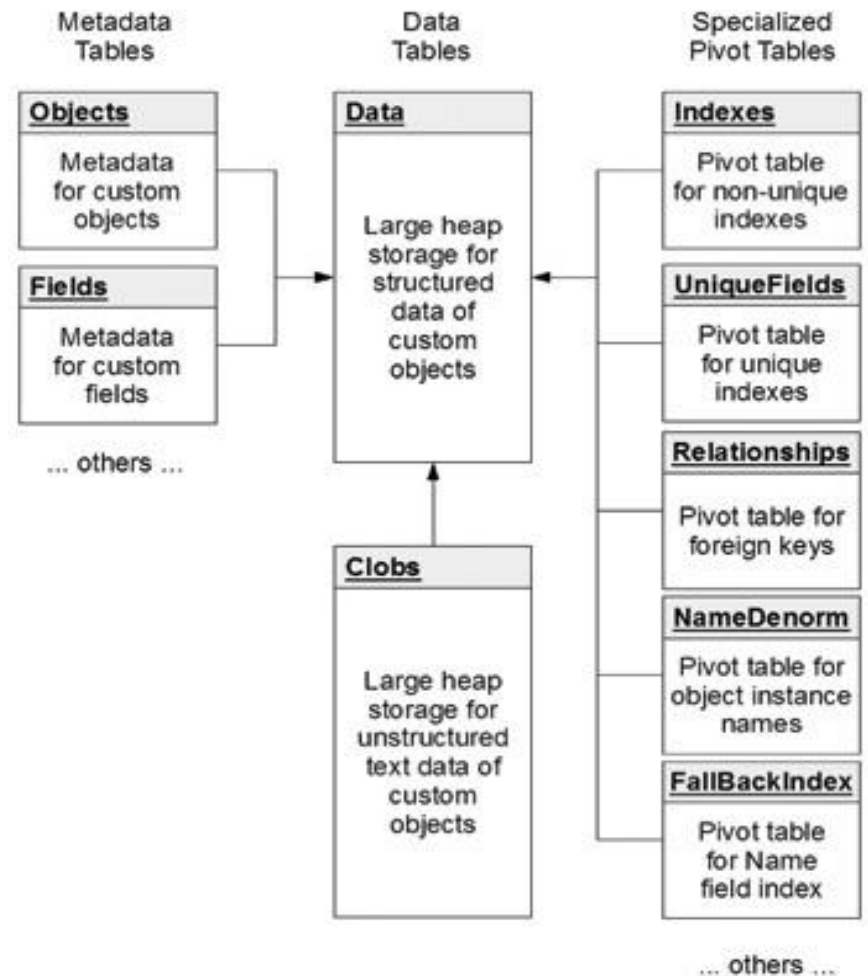


Figure: Force.com data storage model

Q4: Data Storage Model of Force.com

Metadata table – Objects Metadata table

- Two types of metadata tables to store user's customizations
- Objects Metadata table
 - Used to store “*information*” (not the corresponding data values within the object), about the **user-defined objects**.
 - Stored object information is annotated with
 - **OrgID** - A unique identifier assigned to each unique end-user
 - **ObjID** - Force.com assigned unique identifier to each user-created object
 - **ObjectName** – The name of the object
 - Each new user-defined class will have an associated row in this table

<u>Objects</u>
ObjID
OrgID
ObjName

Figure: Metadata table for the storing information on custom classes

Q4: Data Storage Model of Force.com

Metadata tables – Fields Metadata table

- Fields Metadata table
 - Used to store “*information*” (not the corresponding data values) about the **user-defined fields within an object**.
 - Stored field information is annotated with:
 - **OrgID** - The identification of the end-user who owns the object encompassing the field
 - **ObjID** - The identification of the object to whom the field belongs
 - **FieldName** – The name of the field
 - **IsIndexed** – Boolean value to indicate if the field requires indexing
 - **FieldNum** – Value to determine the position of the field in the object relative to other fields
 - Each new user-defined field within an object will have an associated row in this table

<u>Fields</u>
FieldD
OrgID
ObjID
FieldName
Datatype
IsIndexed
FieldNum

Figure: Metadata table for the fields within an object

Q4: Data Storage Model of Force.com - Data table

- Data table used to store the application-specific data
- The stored data is annotated by the following identifiers:
 - **OrgID**, **ObjID**, **ObjName**, and **Global Unique Identifier** (assigned by Force.com to the corresponding data)
- The type of the attributes in the data table is Custom/Flex, to accommodate any structured data types supported by the Force.com platform
- Each new user-defined object will have an associated row in this table

Data
GUID
OrgID
ObjID
Name
Value0
Value1
Value2
...
Value500

Figure: Data table for the fields within an object 16

Q4: Data Storage Model of Force.com

CLOBs tables

- Force.com platform supports data storage as “Text Area (Long)” up to 32, 768 characters. This is used to store arbitrary and unstructured text.
- If an object has “Text Area (Long)” as one of its fields, then the corresponding data is stored in the CLOB (Character Large Objects) table




Figure: Text Area (Long) fields are stored in CLOB table

Q4: Data Storage Model of Force.com

Pivot tables

- The external search engine indexes user's data and stores its summary (based on different criteria), in the pivot tables .
- **Indexes Pivot Table**
 - This table indexes strongly typed data types with in the Force.com platform to enable fast searching on these data types

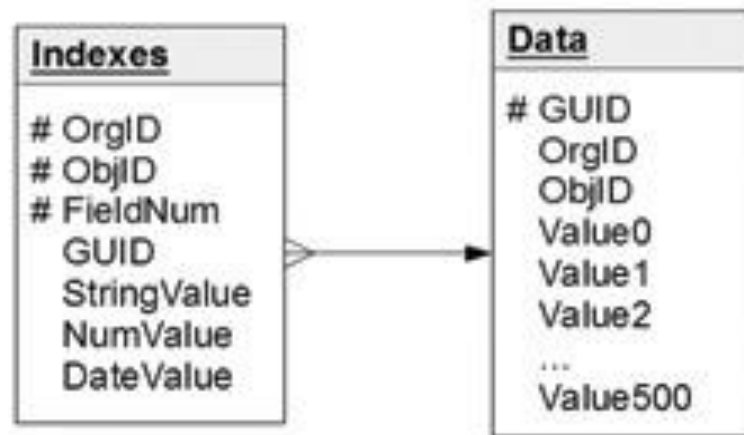
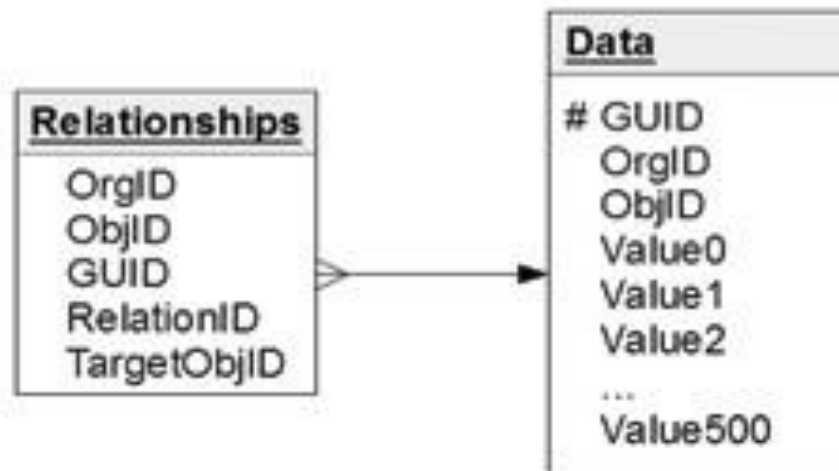


Figure: The index pivot table stores a pivot of the indexed object field

Q4: Data Storage Model of Force.com

Pivot tables

- **UniqueFields Pivot Table**
 - This table is used to enforce the uniqueness of a user-defined field (if defined as a **unique field** within an object by the user)
- **Relationships Pivot Table**
 - This table stores a summary of all the relationships between custom objects;
 - Custom-relationships cannot be defined (at this stage) in Force.com



Q4: Data Storage Model of Force.com

Pivot tables

- Fallback Index Table
 - This index table is used in case the external search engine fails or is overloaded;
 - Stores the names of all the objects to enable a “fall-back” search operation to be performed.
- History Tracking Table
 - This table stores information about changes (field changes, object changes ..etc) carried out by a given user (indexed by OrgID) to enable tracking;
 - Useful in carrying out audit trails.

Q4: Data Storage Model of Force.com

Partitioning of Data and Metadata

- All the data structures (metadata, pivot table structures and data tables) used in the Force.com platform are partitioned by **OrgID** using **native database partitioning mechanism**
- Using the multi-tenant query optimizer, a given query from a given user will only be directed at the **corresponding partitioned data structures**.