

lecture 4: Relational Model

Modern Database Management


11th Edition, International Edition

Chapter 4: Logical Database Design and the Relational Model

Jeffrey A. Hoffer, V. Ramesh, Heikki Topi

Innovation in practice
eng.uts.edu.au • it.uts.edu.au

UTS CRICOS PROVIDER CODE: 00099F

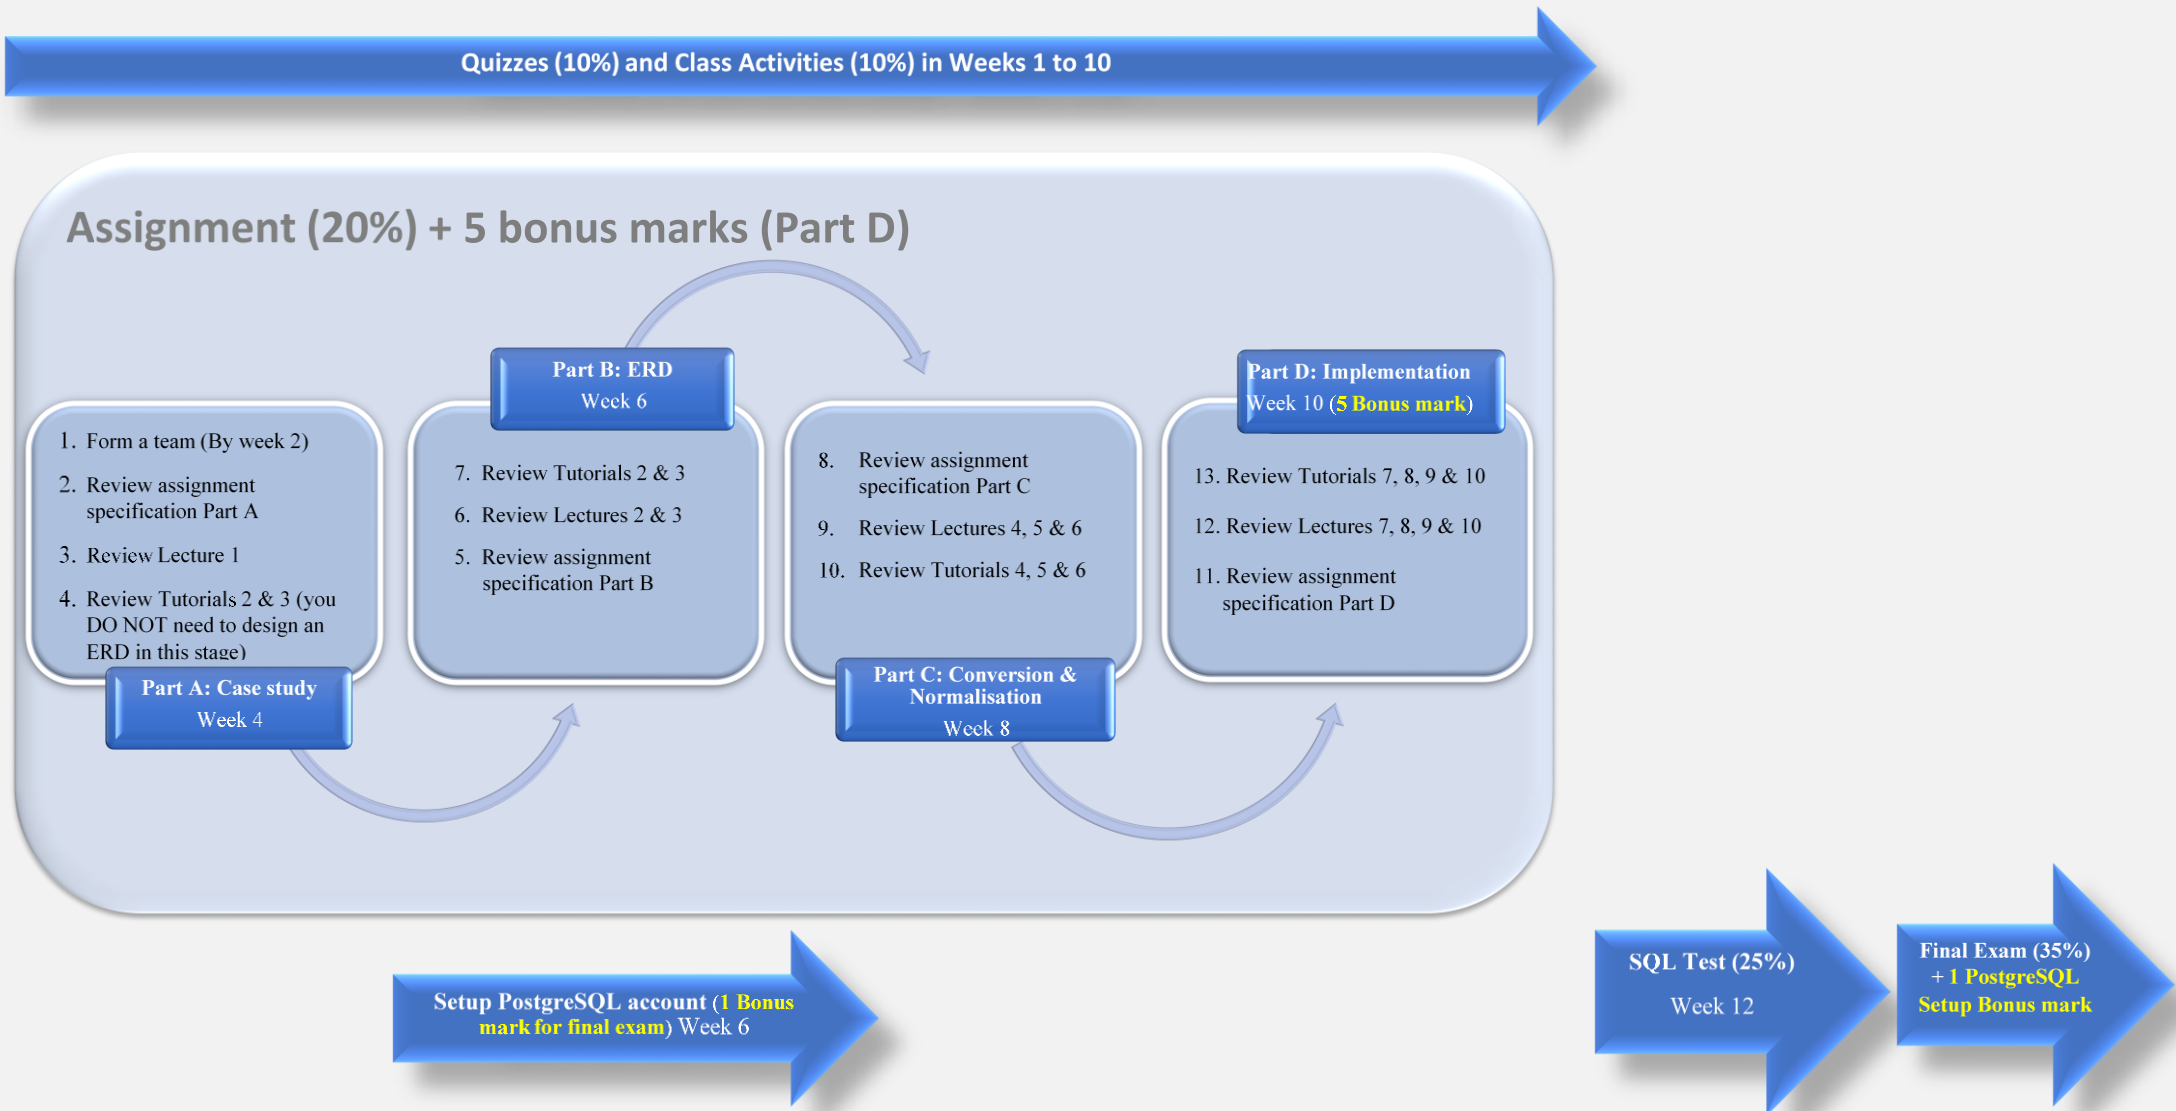


Participations and Discussions

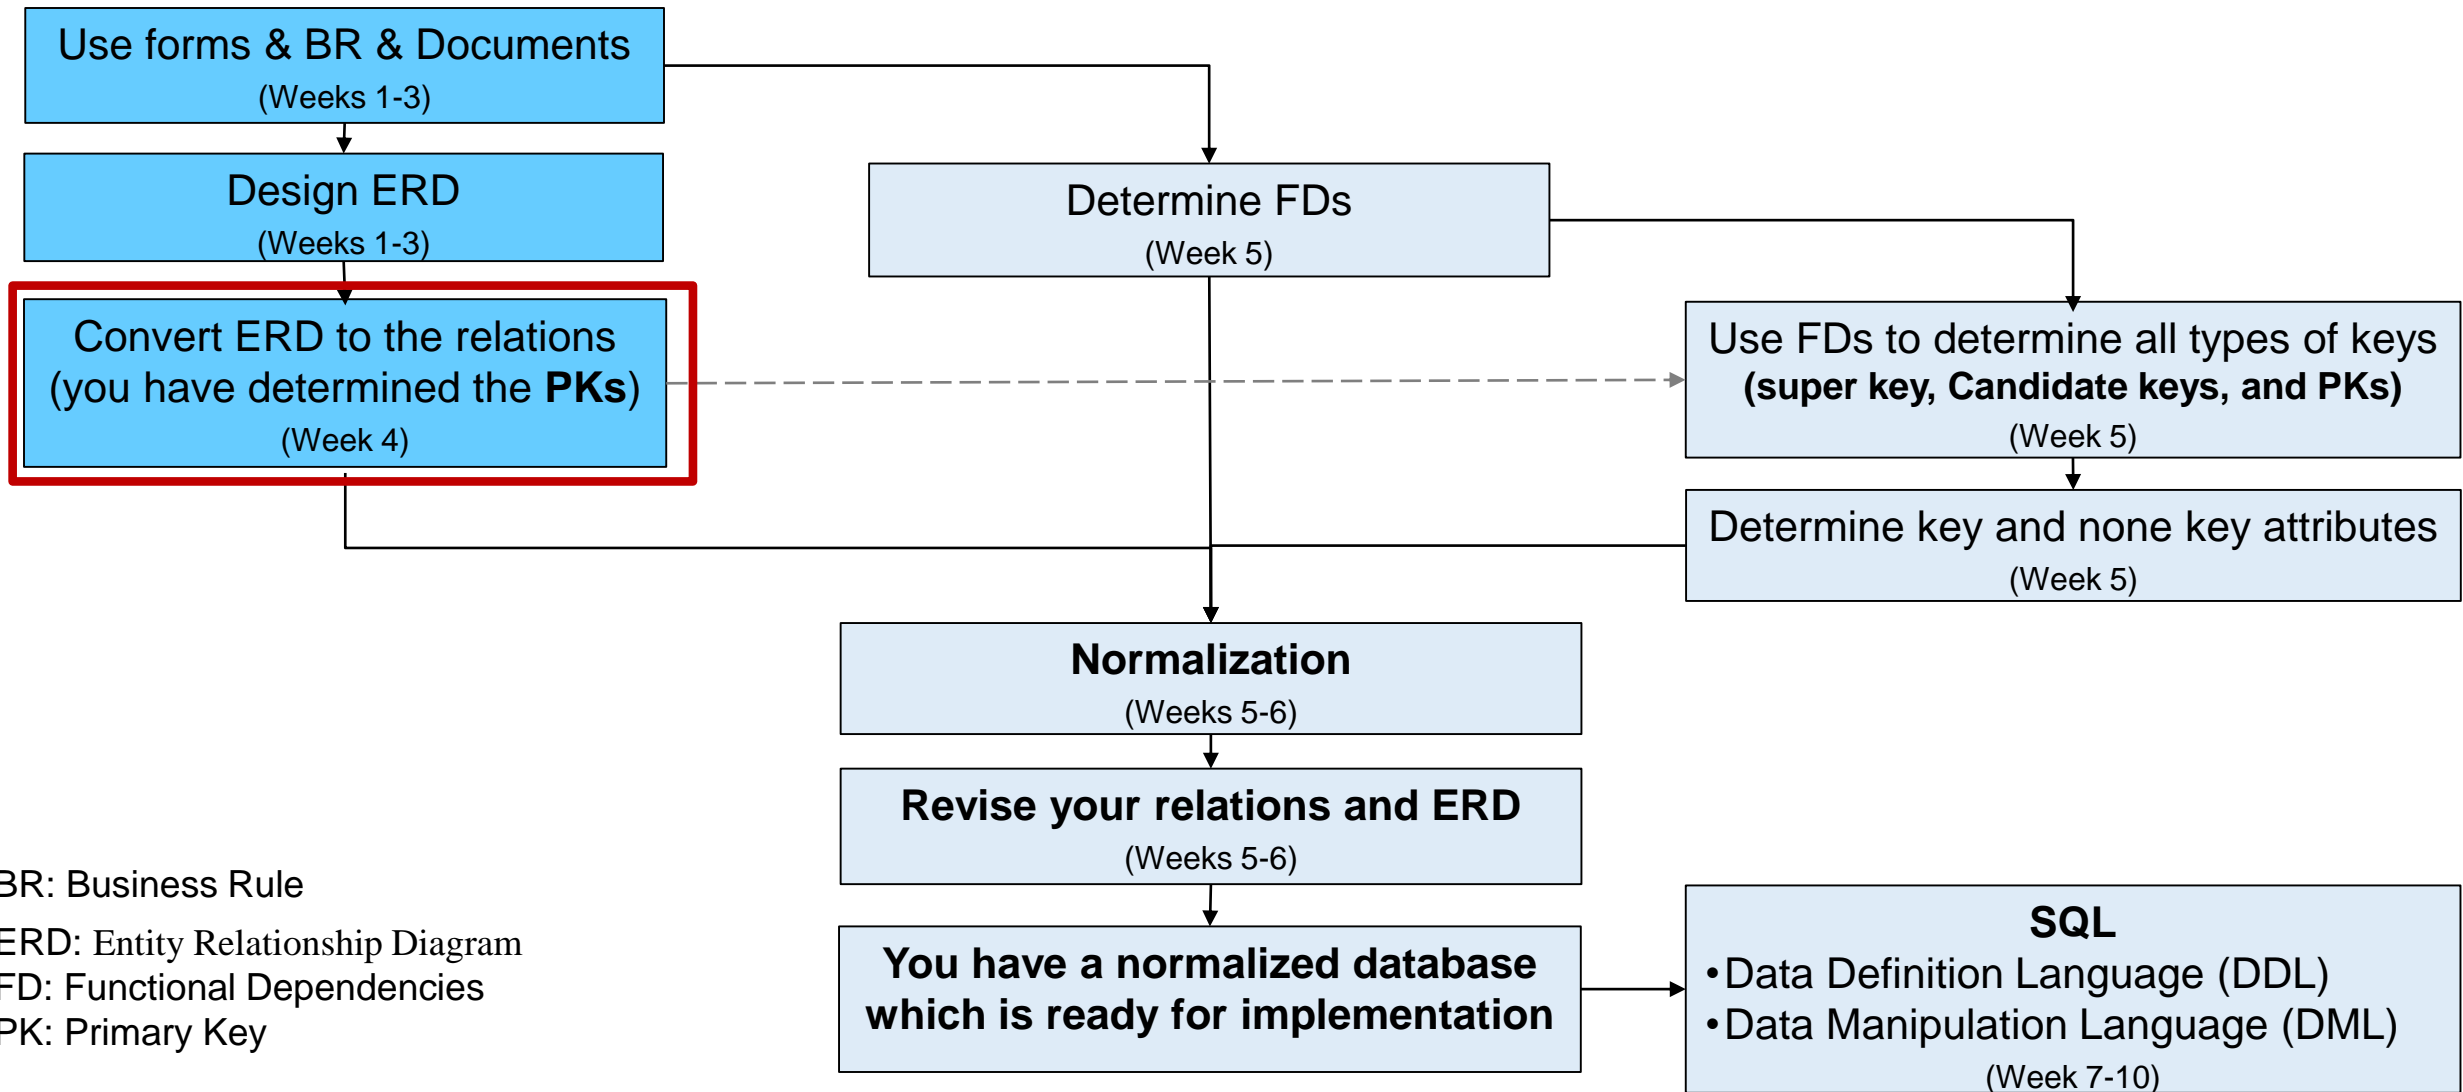
If you have any question and you don't want to share it now,
send it to us via **UTSOnline/Discussion Board**.

However, it is better to speak out 😊

Assessment Chart and Knowledge Guideline



Subject Flowchart



DF Learning Plan

Description: we will have collaborative lecture at the beginning of the class. You need to do some tasks during the lecture as part of your class activities. Then you will do a quiz of what you have learned, then the tutorial will start. you will work in groups during the class.

Please be aware that the lecture slides with Blue title are designed for your self study.

Workshop Timetable:

Activity	Duration	Comments
Lecture	1 hour and 30 minutes	You will have 3 tasks to complete that need to take 20 minutes in total
Rest	10 minutes	Have fun
Review	10 minutes	Please review the review questions and ask your questions if you have any
Tutorial	1 hour	Have even more fun :D (you have two tasks, and need to take be completed in 40 minutes plus 20 minutes for tutors to provide you the solution)
Quiz (Open Book)	5 minutes	On today's content. Will be run before or after the tutorial. Do your best ;)
Leave the class	5 minutes	Don't forget to review what you have learn in this class, and check the information that is provided on UTSOnline/Learning Material/Week 4

Subject Overview

➤ Design Entity Relationship Diagram (ERD)

- Week 1: Data Modelling I (Conceptual Level)
- Week 2: Data Modelling II (Conceptual Level)
- Week 3: Data Modelling III (Conceptual Level)
- **Week 4: Convert ERD to Relations (Logical Level)**
- Week 5: Functional Dependencies
- Week 5: Normalization I
- Week 6: Normalization II

➤ Data manipulation

- Week 7: Simple Query
- Week 8: Multiple Table Queries
- Week 9: Subquery
- Week 10: Correlated Subquery

Objectives

1. Components of relational model

2. Relations

2.1. Correspondence with E-R Model

2.2. Key Fields

2.3. Integrity Constraints

2.3.1. Domain Constraints

2.3.2. Entity Integrity

2.3.3. Referential Integrity

3. Transforming EER Diagrams into Relations

3.1. Mapping Regular Entities to Relations (with simple, composite, and multivalued attributes)

3.2. Mapping Weak Entities

3.3. Mapping Binary Relationships (1:M, M:N, 1:1)

3.4. Mapping Associative Entities

3.5. Mapping Unary Relationships

3.6. Mapping Ternary (and n-ary) Relationships

3.7. Mapping Supertype/Subtype Relationships

1. Components of Relational Model

- **Data structure**
Tables (relations), rows, columns
- **Data manipulation**
Powerful SQL operations for retrieving and modifying data
- **Data integrity**
Mechanisms for implementing business rules that maintain integrity of manipulated data

2. Relation

2. Relation


- A relation is a named, two-dimensional **table** of data.
- A table consists of **rows** (records) and **columns** (attribute or field).
- Requirements for a table to qualify as a relation:
 - It must have a **unique name**.
 - Every attribute value must be **atomic** (not multivalued, not composite) (*More on this in the next lectures*).
 - Every **row** must be **unique** (can't have two rows with exactly the same values for all their fields).
 - **Attributes** (columns) in tables must have **unique names**.
 - The order of the columns must be irrelevant.
 - The order of the rows must be irrelevant.

2.1. Correspondence with E-R Model

- **Relations (tables)** correspond with **entity types** and with many-to-many relationship types.
- **Rows** correspond with **entity instances** and with many-to-many relationship instances.
- **Columns** correspond with **attributes**.

NOTE: The word *relation* (in relational database) is NOT the same as the word *relationship* (in E-R model).

2.2. Key Fields



- Keys are special fields that serve two main purposes:



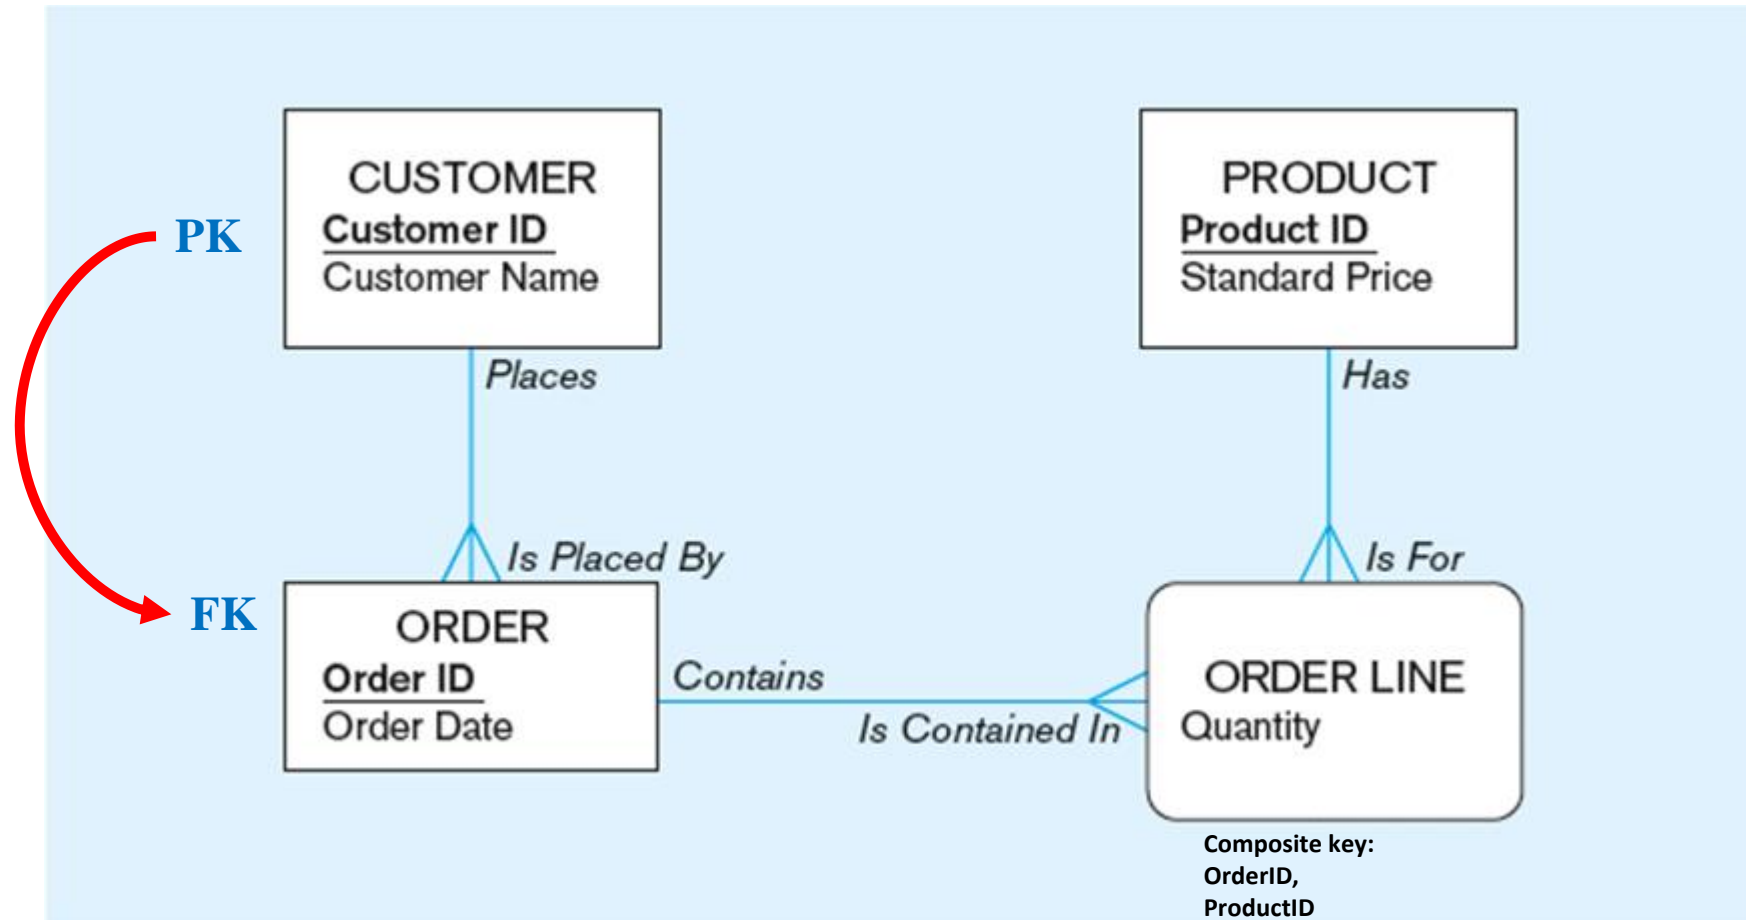
- **Primary keys** are **unique** identifiers of the relation.
 - Examples include employee numbers, social security numbers, etc. This guarantees that all rows are unique.



- **Foreign keys** are identifiers that enable a **dependent** relation (on the many side of a relationship) to refer to its **parent** relation (on the one side of the relationship).

- Keys can be **simple** (a single field) or **composite** (more than one field).
- Keys usually are used as **indexes** to speed up the response to user queries (*more on this in Chapter 5*).

ERD of Pine Valley Furniture Company



Primary key/Foreign key

Business Rule: any order can be related to many products.
(related to OrderLine_T)

PK

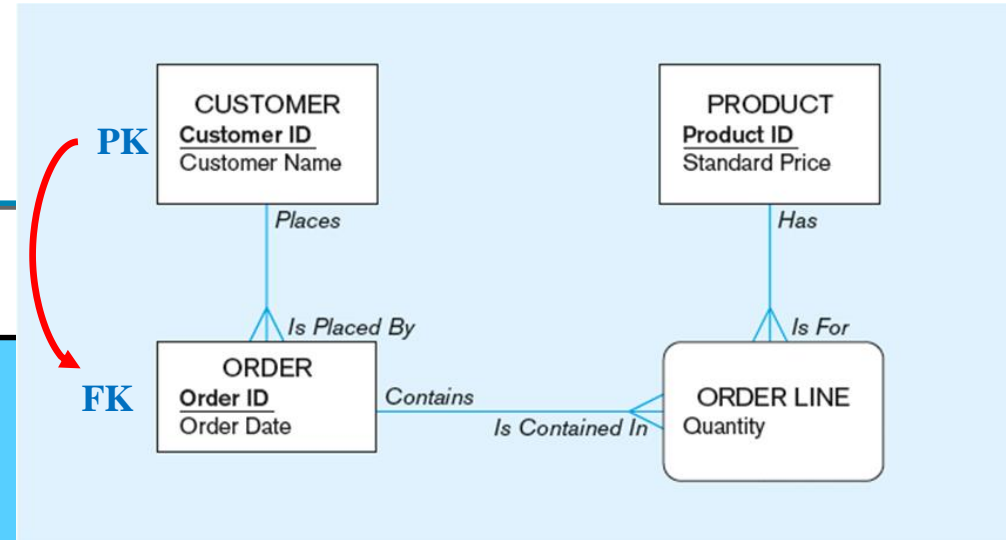
CustomerID	CustomerName
1	Contemporary Casuals
2	Value Furniture
3	Home Furnishings
4	Eastern Furniture
5	Impressions
6	Furniture Gallery
7	Period Furniture
8	California Classics
9	M and H Casual Furniture
10	Seminole Interiors
11	American Euro Lifestyles
12	Battle Creek Furniture
13	Heritage Furnishings
14	Kaneohe Homes
15	Mountain Scenes

OrderID	OrderDate	CustomerID
1001	10/21/2015	1
1002	10/21/2015	8
1003	10/22/2015	15
1004	10/22/2015	5
1005	10/24/2015	3
1006	10/24/2015	2
1007	10/27/2015	11
1008	10/30/2015	12
1009	11/5/2015	4
1010	11/5/2015	1

FK

OrderID	ProductID	OrderedQuantity
1001	1	2
1001	2	2
1001	4	1
1002	3	5
1003	3	3
1004	6	2
1004	8	2
1005	4	4
1006	4	1
1006	5	2
1006	7	2
1007	1	3
1007	2	2
1008	3	3
1008	8	3
1009	4	2
1009	7	3
1010	8	10

ProductID	ProductDescription	ProductPrice
1	End Table	Cherry
2	Coffee Table	Natura
3	Computer Desk	Natura
4	Entertainment Center	Natura
5	Writers Desk	Cherry
6	8-Drawer Desk	White
7	Dining Table	Natura
8	Computer Desk	Walnut



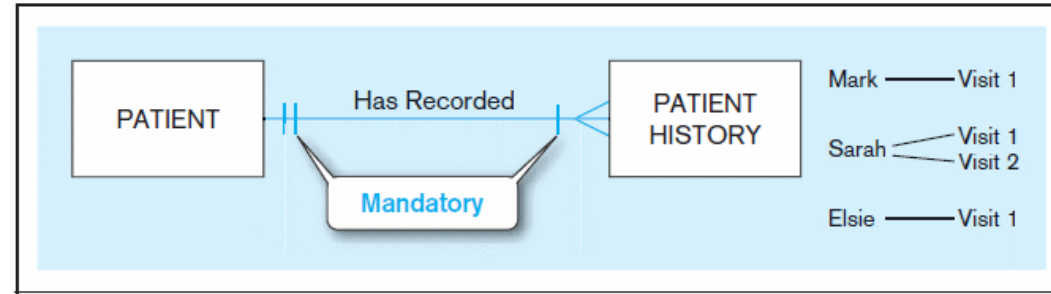
If two relations (tables) have a relationship, then the PK of the parent relation will be a FK in the dependent relation ...

Question: What would be the FK in a relation (table) that need to have a relationship with OrderLine_T?

Class Activity 4.1 (5 minutes)


1. Explain why when there is one-to-many relationship between two entities, **PK** of the entity on the **one side** will be **FK** on the entity **on the many side**. Provide an example to explain this.

Solution to Class Activity 4.1:



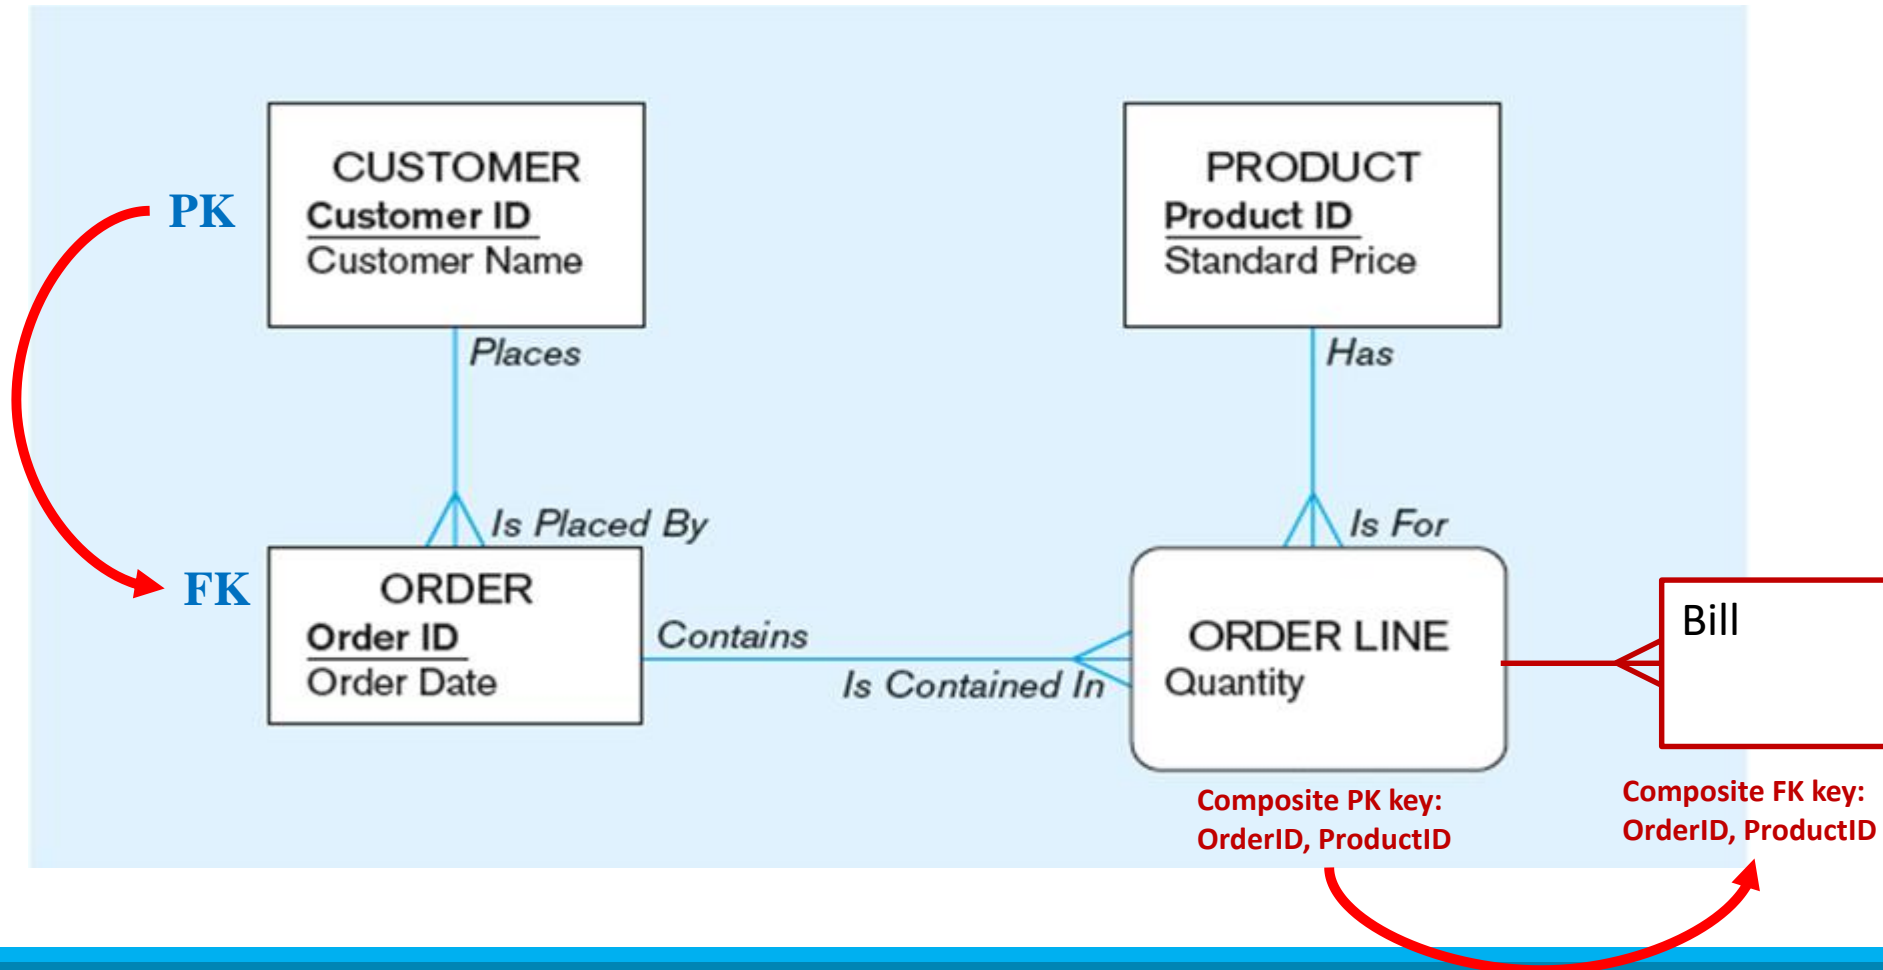
Discussion: Why in each **one-to-many** relationships, **PK of the entity on one side** is **FK of the entity on the many side**? (From Lecture 2)

BR: One instructor **can** teach **many** subjects, but one subject **needs to** be taught by **one** instructor.

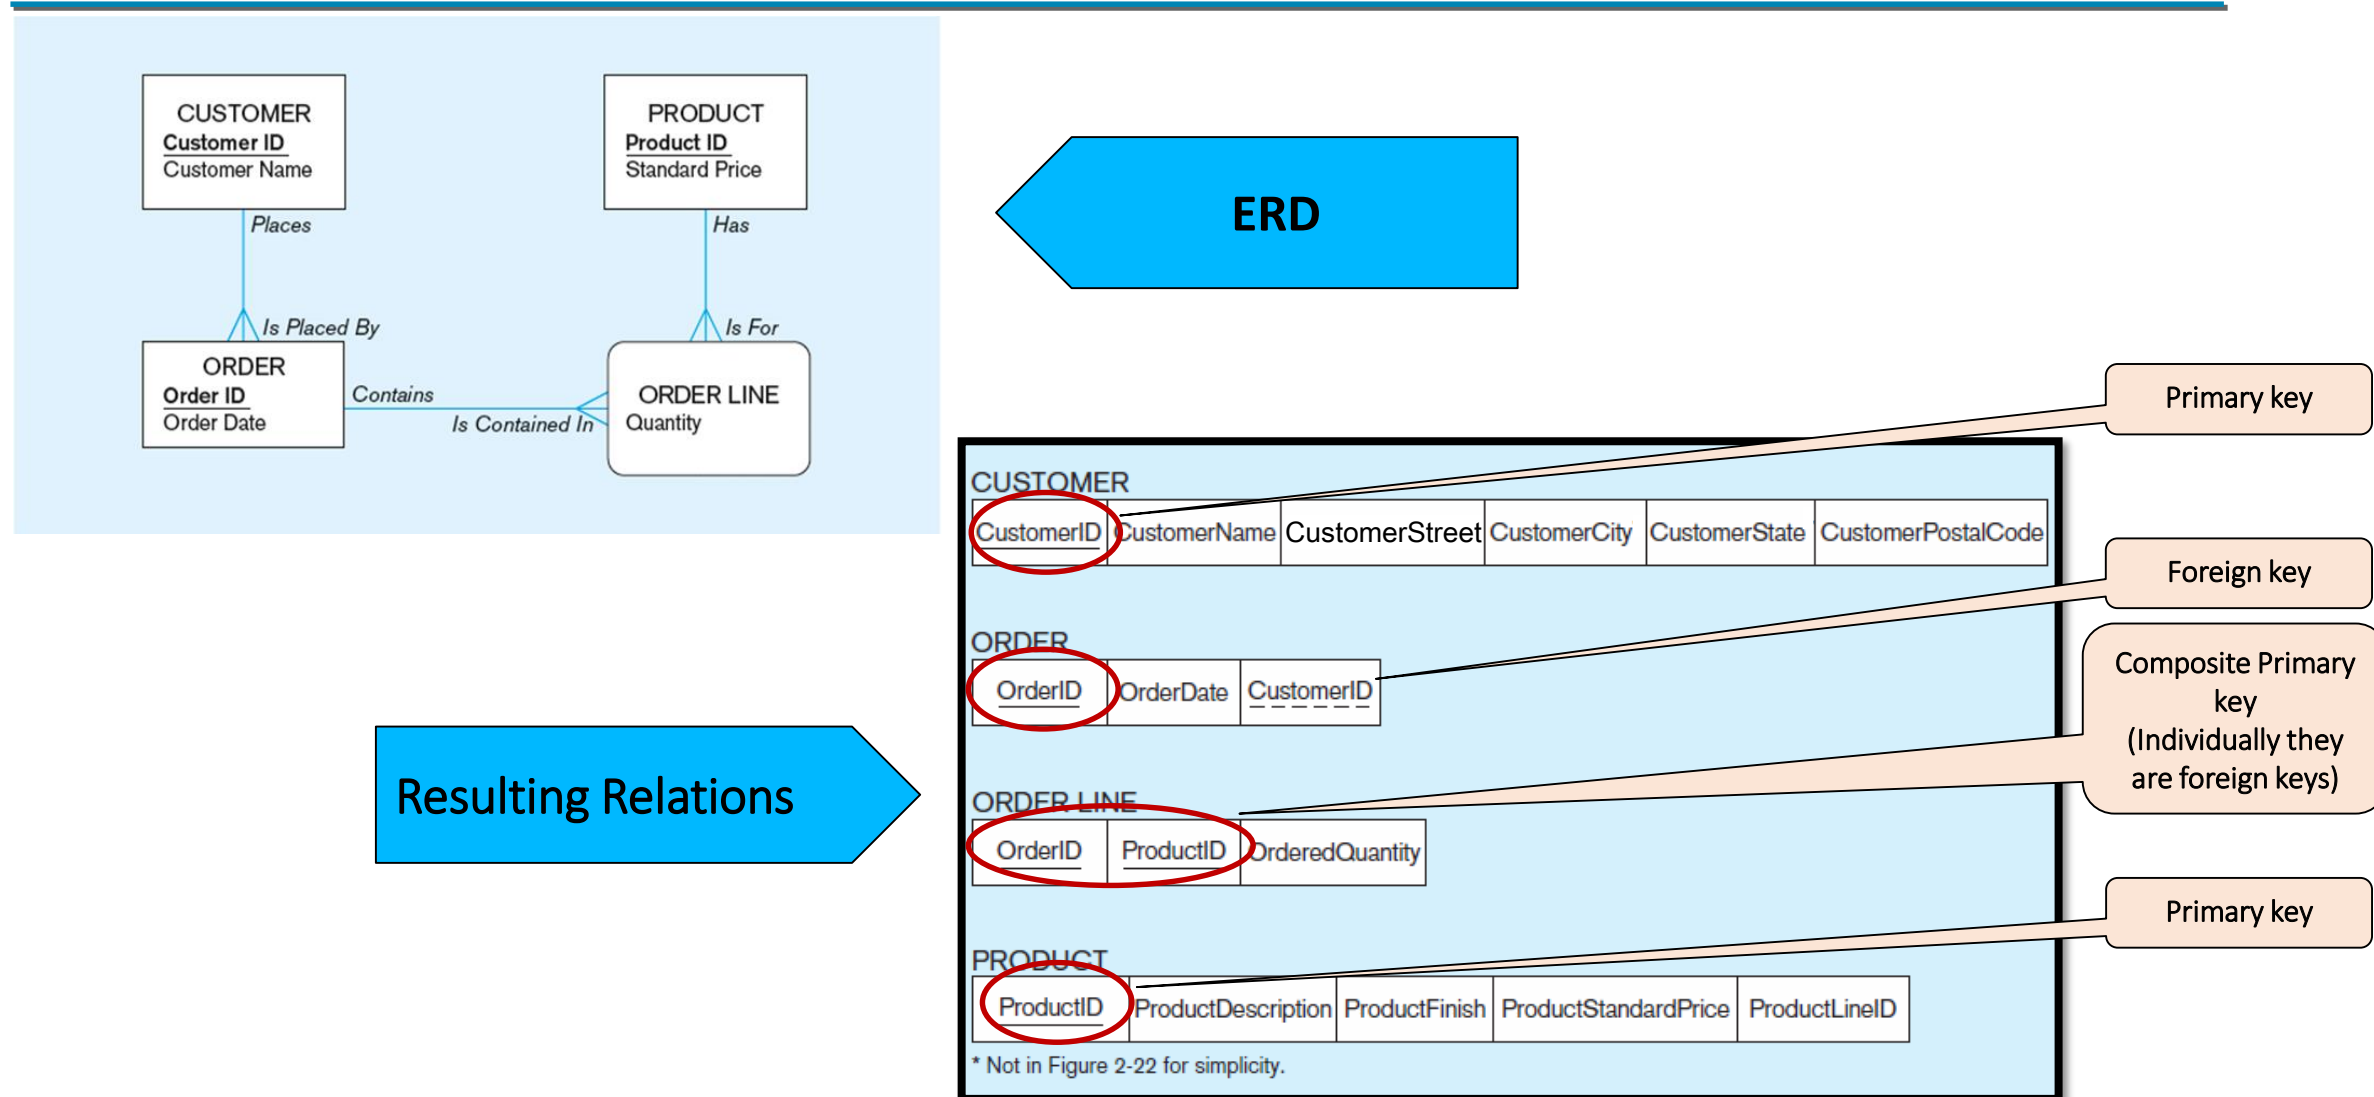


ERD of Pine Valley Furniture Company

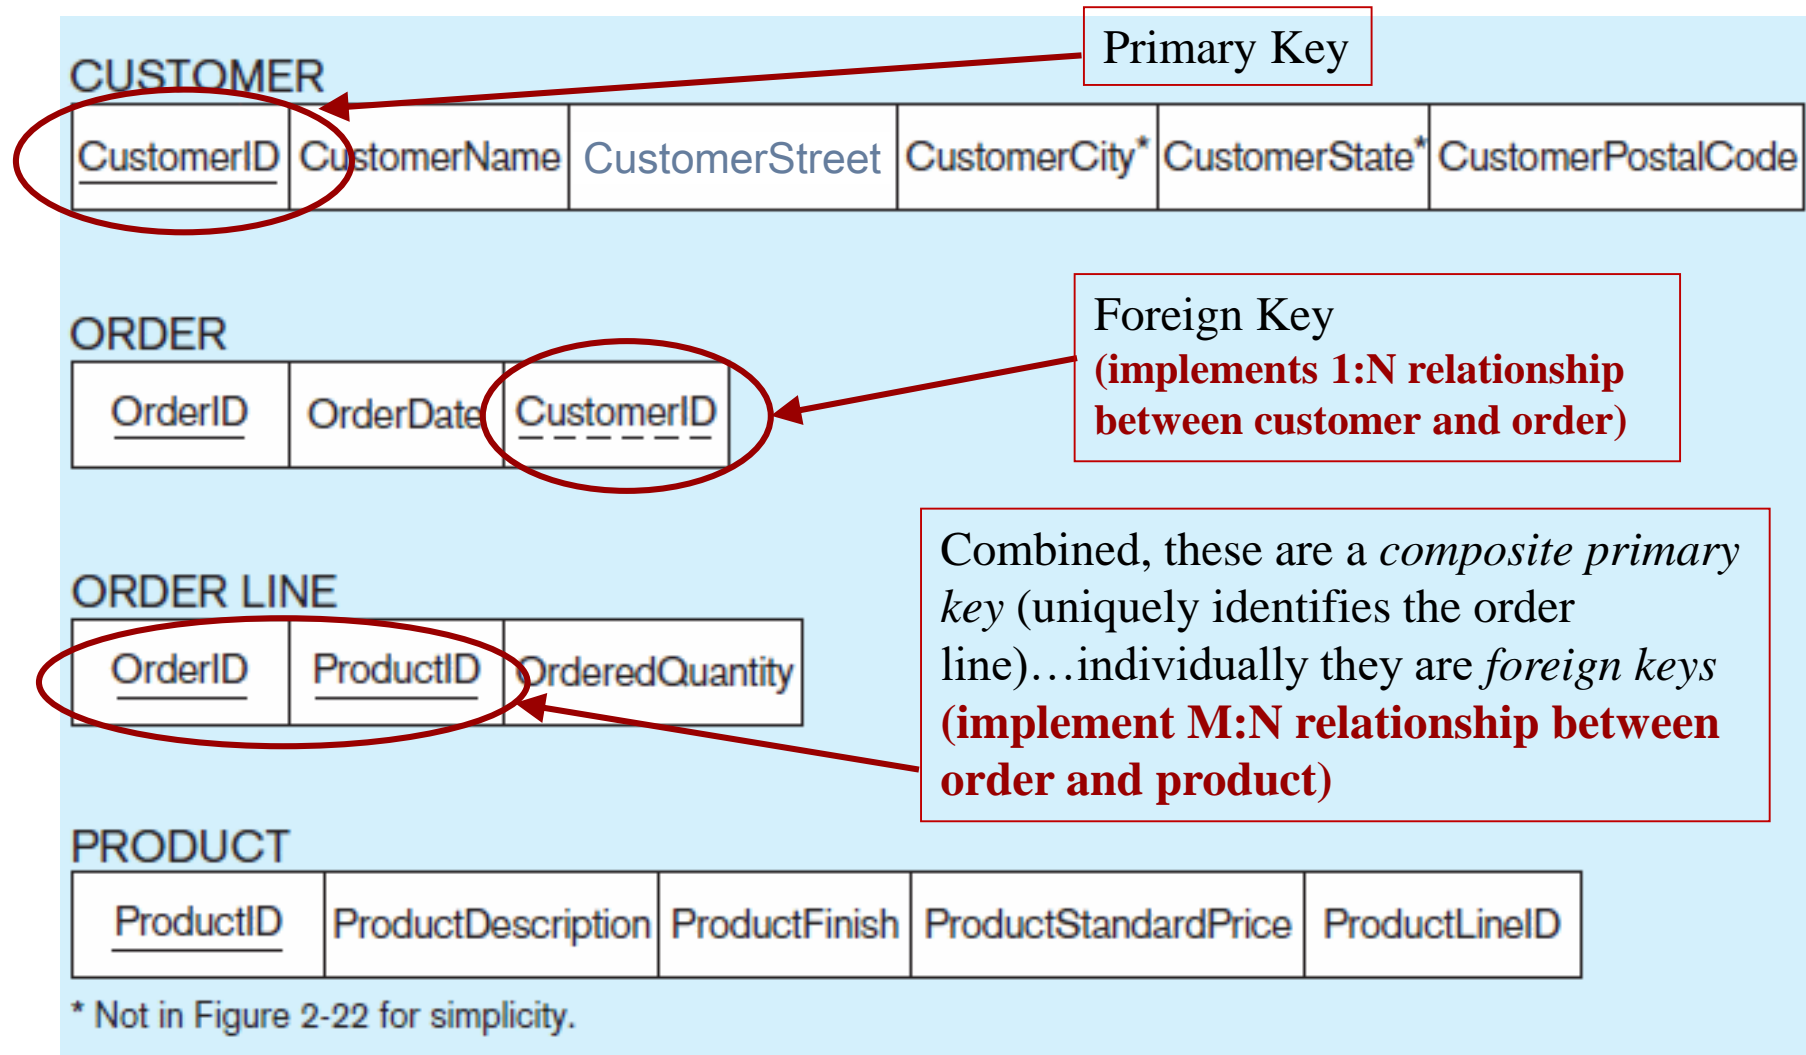
Question: What would be the FK in a relation (table) that need to have a relationship with OrderLine_T?



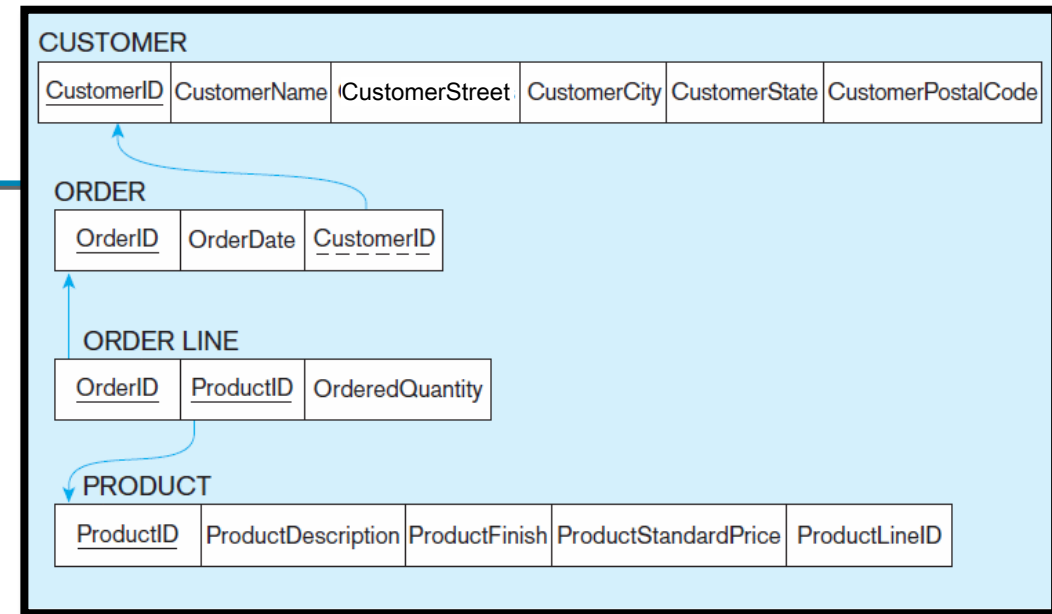
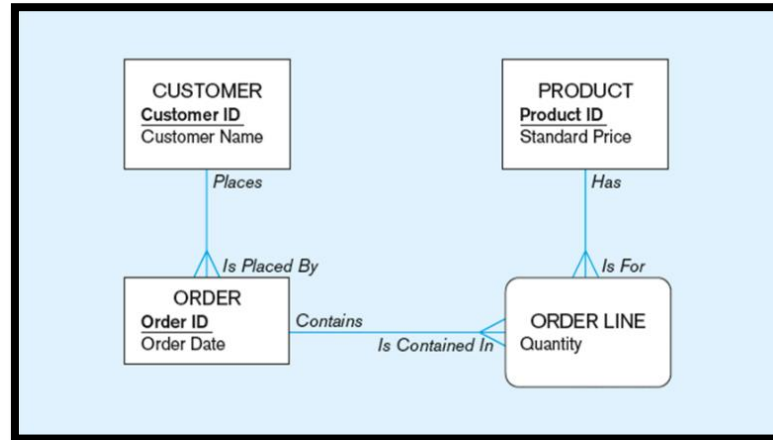
Schema for four relations (Pine Valley Furniture Company)-(Figure 4-3)



Schema for four relations (Pine Valley Furniture Company)-(Figure 4-3)



Schema to relations Methods:



CUSTOMER (CustomerID, CustomerName, CustomerStreet, CustomerCity, CustomerState, CustomerPostalCode)

ORDER (OrderID, OrderDate, CustomerID^{*})

FK (CustomerID) references CUSTOMER

ORDERLINE (OrderID^{*}, ProductID^{*}, OrderQuantity)

FK (OrderID) references ORDER

FK (ProductID) references PRODUCT

PRODUCT (ProductID, ProductDescription, ProductFinish, ProductStandardPrice, ProductLineID)

Where we are ... review the path 😊

- ❑ We know about the relations
- ❑ We have abstract information about converting ERD to the Relations.
- ❑ Now we need to know which constraints need to be considered when we do the conversion and why.

2.3. Integrity Constraints

Integrity Constraints are applied to facilitate maintaining the accuracy and integrity of data in the database. The major types of integrity constraint are:

2.3.1. Domain Constraints

- Allowable values for an attribute (See Table 4-1)

2.3.2. Entity Integrity

- No primary key attribute may be **null**. All primary key fields **MUST** have data.

2.3.3. Referential Integrity

- states that **any foreign key** value (on the relation of the many side) **MUST match a primary key** value in the relation of the one side.

Referential Integrity rule is used to maintain the consistency among rows between the two tables.

2.3.1. Domain Constraints

Domain Constraints Allowable values for an attribute (See Table 4-1)

TABLE 4-1 Domain Definitions for INVOICE Attributes			
Attribute	Domain Name	Description	Domain
CustomerID	Customer IDs	Set of all possible customer IDs	character: size 5
CustomerName	Customer Names	Set of all possible customer names	character: size 25
CustomerAddress	Customer Addresses	Set of all possible customer addresses	character: size 30
CustomerCity	Cities	Set of all possible cities	character: size 20
CustomerState	States	Set of all possible states	character: size 2
CustomerPostalCode	Postal Codes	Set of all possible postal zip codes	character: size 10
OrderID	Order IDs	Set of all possible order IDs	character: size 5
OrderDate	Order Dates	Set of all possible order dates	date: format mm/dd/yy
ProductID	Product IDs	Set of all possible product IDs	character: size 5
ProductDescription	Product Descriptions	Set of all possible product descriptions	character: size 25
ProductFinish	Product Finishes	Set of all possible product finishes	character: size 15
ProductStandardPrice	Unit Prices	Set of all possible unit prices	monetary: 6 digits
ProductLineID	Product Line IDs	Set of all possible product line IDs	integer: 3 digits
OrderedQuantity	Quantities	Set of all possible ordered quantities	integer: 3 digits

Domain definitions enforce domain integrity constraints.

2.3.3. Referential Integrity

Referential Integrity states that **any foreign key** value (on the relation of the many side) **MUST match a primary key** value in the relation of the one side.

Referential Integrity rule is used to maintain the consistency among rows between the two tables.

Example of integrity constraint: Delete Rules

Restrict—don't allow delete of “parent” side if related rows exist in “dependent” side

Cascade—automatically delete “dependent” side rows that correspond with the “parent” side row to be deleted

Set-to-Null—set the foreign key in the dependent side to null if deleting from the parent side

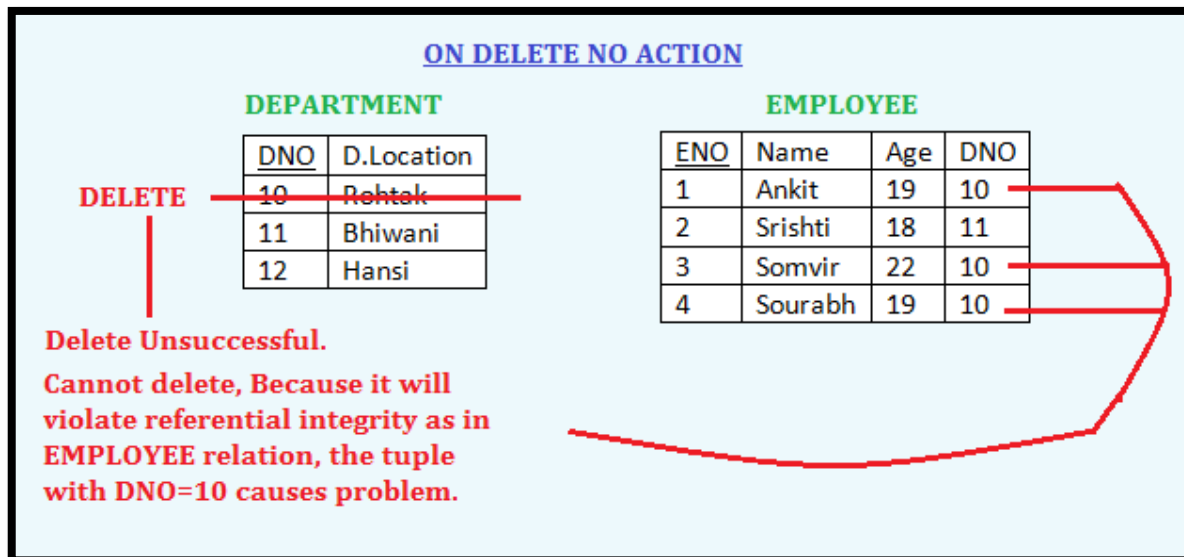
→ The foreign key can be null

→ **Set-to-Null** is not allowed for weak and associated entities

2.3.3. Referential Integrity: Restrict

Restrict: don't allow delete of "parent" side if related rows exist in "dependent" side

➤ Delete of "parent" side



➤ Delete of "dependent" side

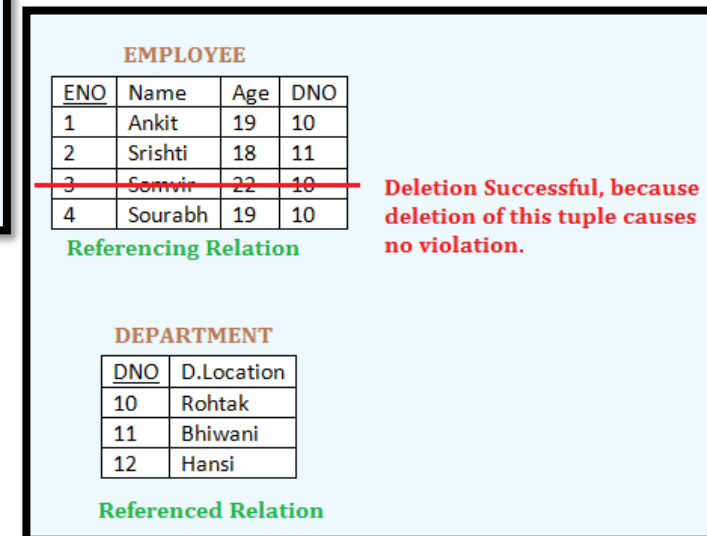


Photo Reference: <http://www.edugrabs.com>

2.3.3. Referential Integrity: Cascade

Cascade—automatically delete “dependent” side rows that correspond with the “parent” side row to be deleted

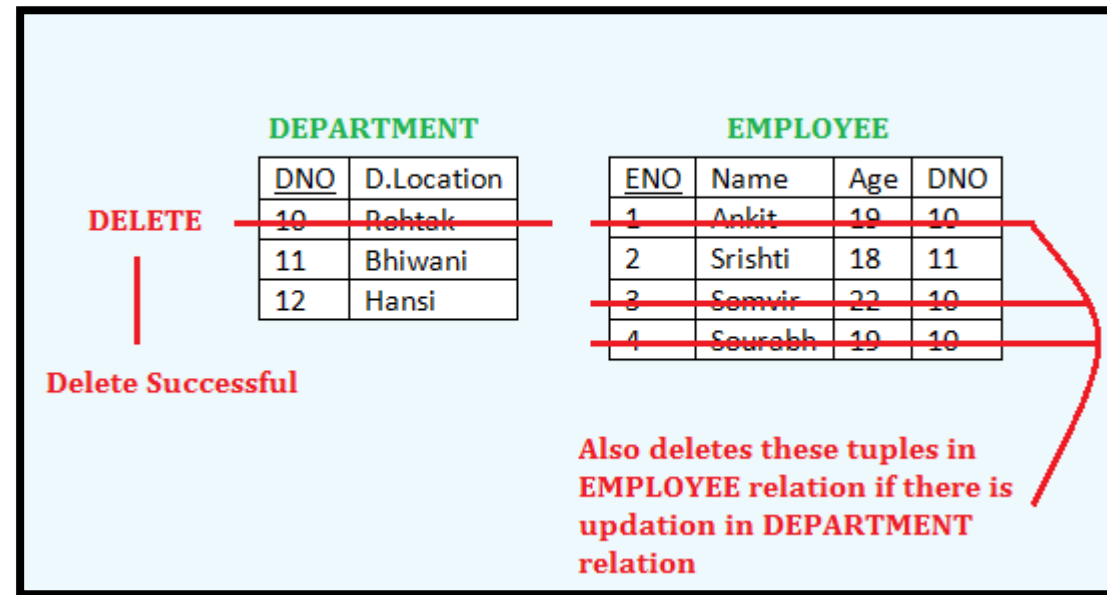


Photo Reference: <http://www.edugrabs.com>

2.3.3. Referential Integrity: Set-to-Null

Set-to-Null: set the foreign key in the dependent side to null if deleting from the parent side.

Notes:

- The foreign key can be null
- **Set-to-Null** is **not allowed** for weak and associated entities (where FK are part of the key).
- **Set-to-Null** is **not allowed** when is related to a mandatory cardinality.

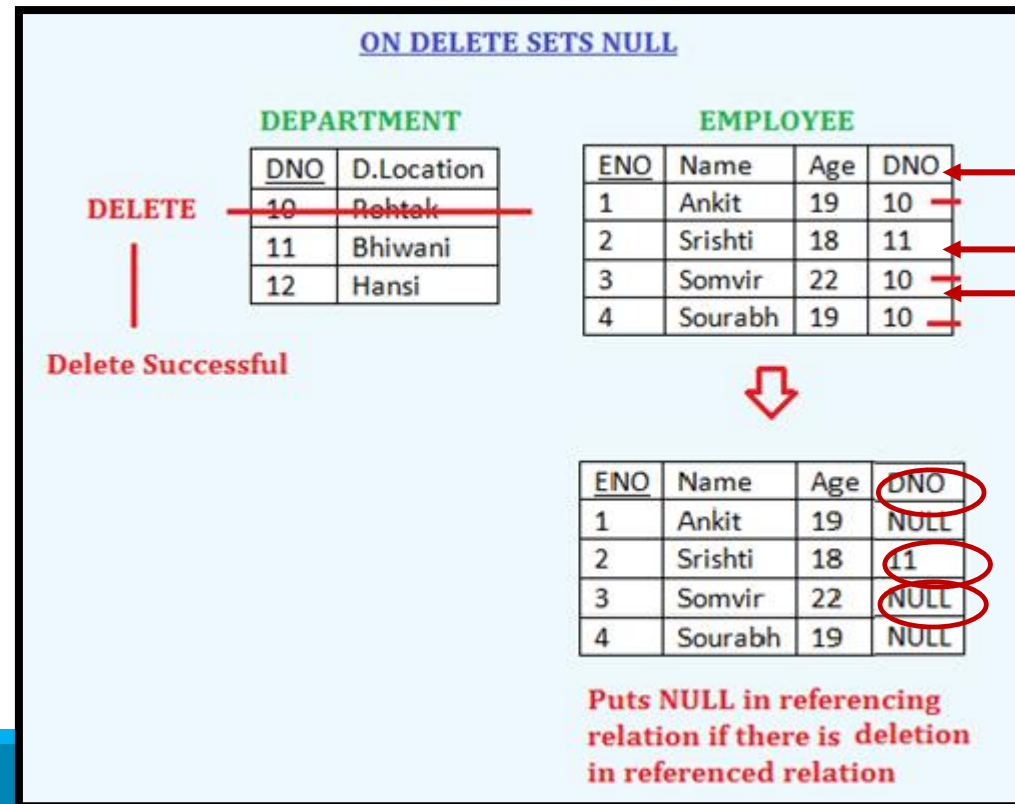
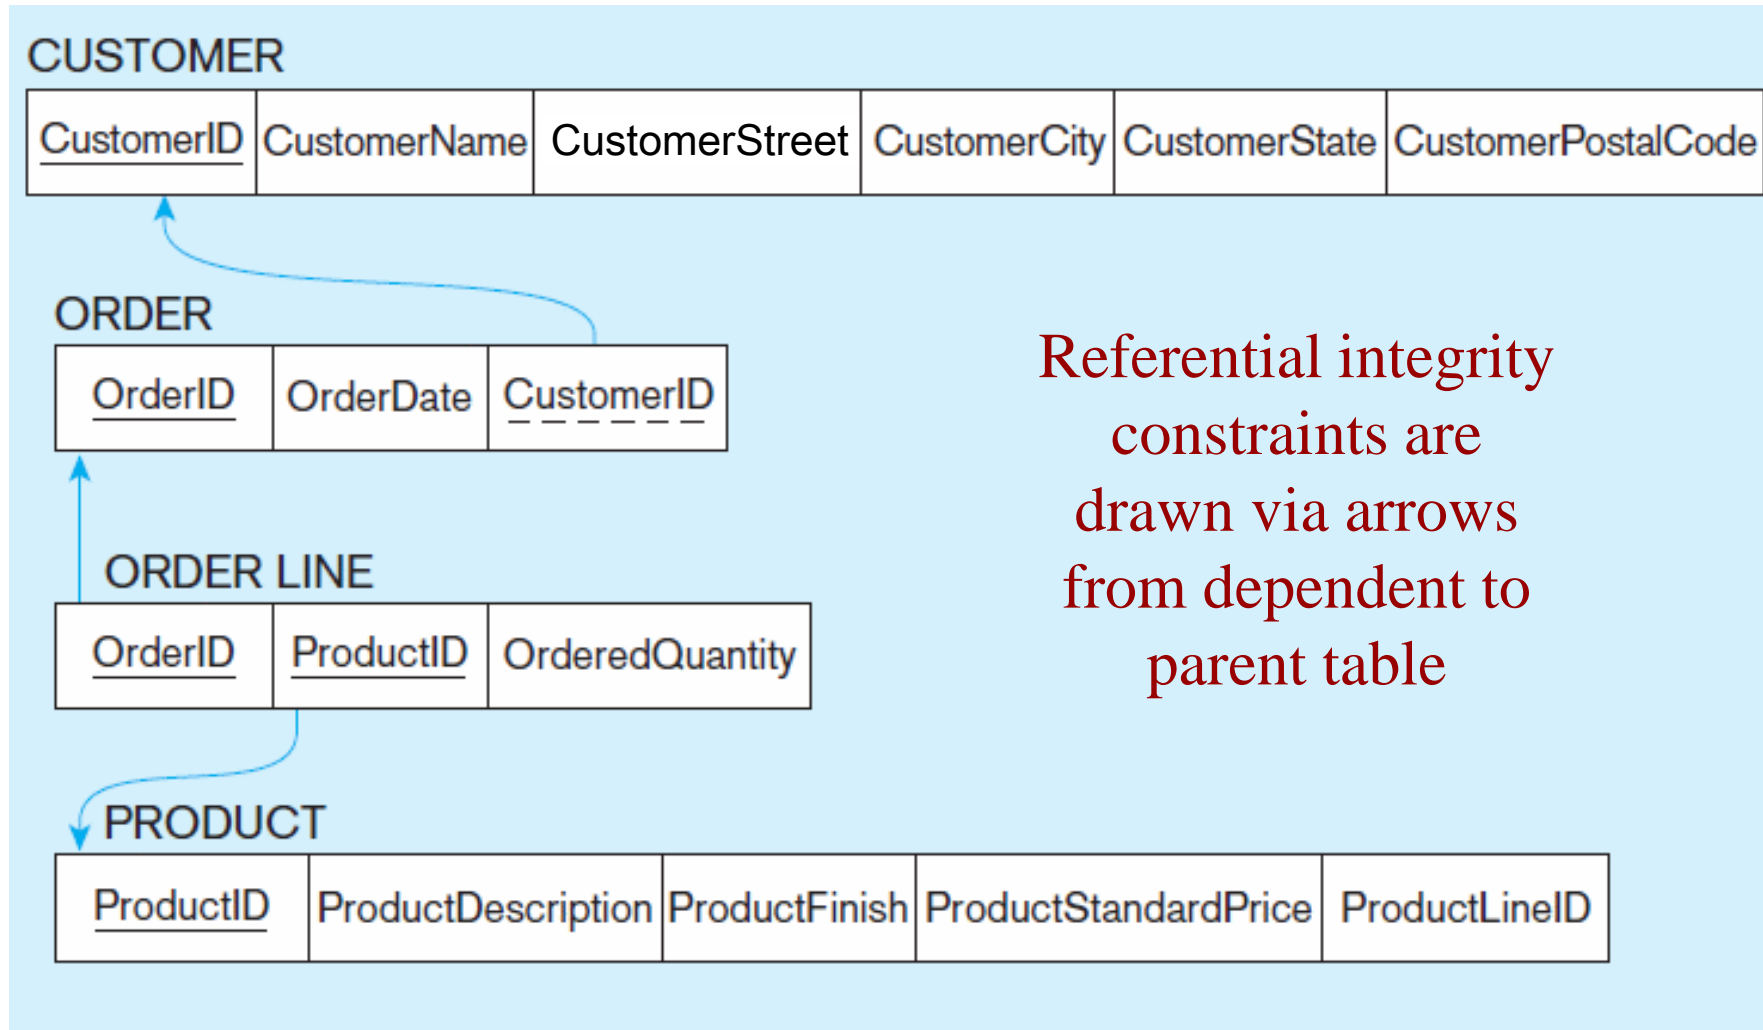


Photo Reference: <http://www.edugrabs.com>

Referential integrity constraints (Pine Valley Furniture)- (Figure 4-5)



SQL table definitions (Figure 4-6)

```
CREATE TABLE Customer_T
  (CustomerID          NUMBER(11,0)    NOT NULL,
   CustomerName        VARCHAR2(25)    NOT NULL,
   CustomerStreet      VARCHAR2(30),
   CustomerCity        VARCHAR2(20),
   CustomerState       CHAR(2),
   CustomerPostalCode  VARCHAR2(9),
   CONSTRAINT Customer_PK PRIMARY KEY (CustomerID));

CREATE TABLE Order_T
  (OrderID            NUMBER(11,0)    NOT NULL,
   OrderDate          DATE DEFAULT SYSDATE,
   CustomerID         NUMBER(11,0),
   CONSTRAINT Order_PK PRIMARY KEY (OrderID),
   CONSTRAINT Order_FK FOREIGN KEY (CustomerID) REFERENCES Customer_T (CustomerID));

CREATE TABLE Product_T
  (ProductID          NUMBER(11,0)    NOT NULL,
   ProductDescription  VARCHAR2(50),
   ProductFinish      VARCHAR2(20),
   ProductStandardPrice DECIMAL(6,2),
   ProductLineID      NUMBER(11,0),
   CONSTRAINT Product_PK PRIMARY KEY (ProductID));

CREATE TABLE OrderLine_T
  (OrderID            NUMBER(11,0)    NOT NULL,
   ProductID          NUMBER(11,0)    NOT NULL,
   OrderedQuantity    NUMBER(11,0),
   CONSTRAINT OrderLine_PK PRIMARY KEY (OrderID, ProductID),
   CONSTRAINT OrderLine_FK1 FOREIGN KEY (OrderID) REFERENCES Order_T (OrderID),
   CONSTRAINT OrderLine_FK2 FOREIGN KEY (ProductID) REFERENCES Product_T (ProductID));
```

Referential integrity constraints are implemented with foreign key to primary key references.

Note: Review this slide after Lecture 7 when you learn about DDL 😊

Where we are ... review the path 😊

- ❑ We know about the relations
- ❑ We have abstract information about converting ERD to the Relations.
- ❑ We know which constraints (integrity constraints) need to be considered when we do the conversion and why.
- ❑ Now we need to know how to convert different type of attributes and entities to the relations.

3. Transforming ERD into Relations

3. Transforming ERD into Relations

- 3.1. Mapping **Regular** Entities to Relations
- 3.2. Mapping **Weak** Entities
- 3.3. Mapping **Binary** Relationships
- 3.4. Mapping **Associative** Entities
- 3.5. Mapping **Unary** Relationships
- 3.6. Mapping **Ternary** (and n-ary) Relationships
- 3.7. Mapping **Supertype/Subtype** Relationships

3.1. Mapping Regular Entities to Relations

3.1.1. Simple attributes: E-R attributes map directly onto the relation

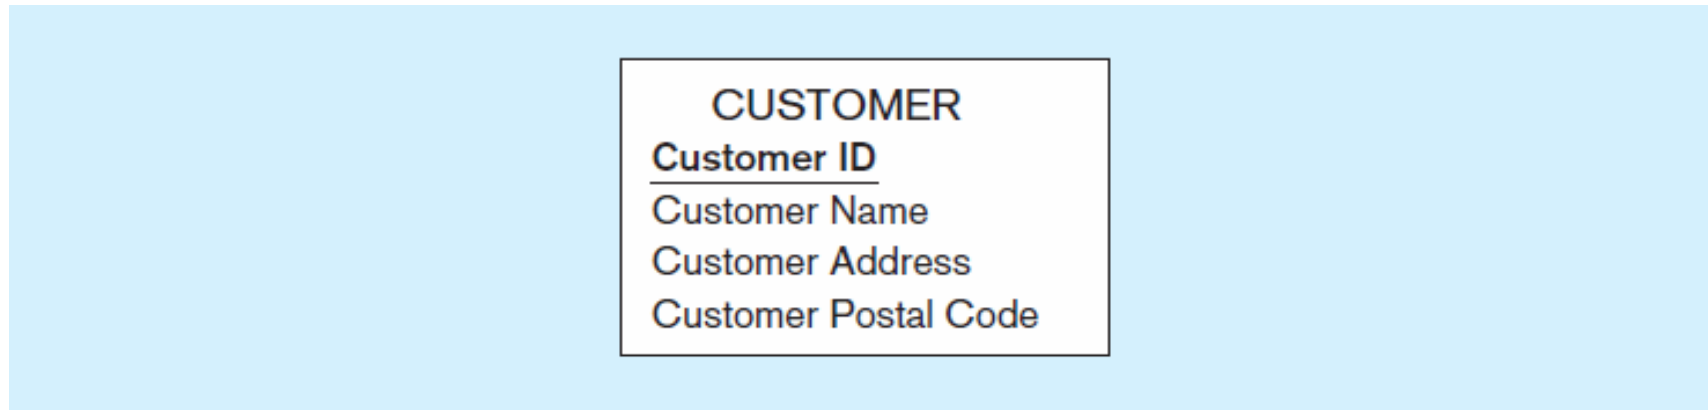
3.1.2. Composite attributes: Use only their simple component attributes

3.1.3. Multivalued Attribute: Becomes a **separate relation** with a foreign key taken from the superior entity


3.1.1. Mapping a regular entity (Figure 4-8) with simple attribute

3.1.1. Simple attributes: E-R attributes map directly onto the relation

(a) CUSTOMER entity type with simple attributes



(b) CUSTOMER relation



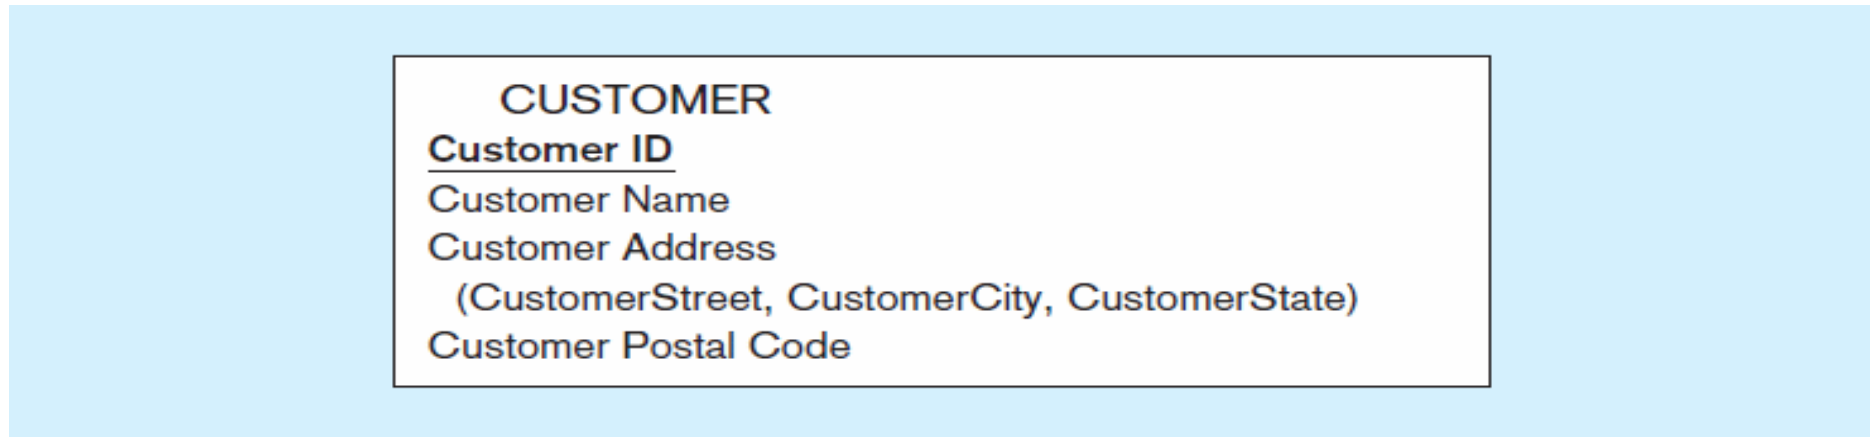
or

CUSTOMER (CustomerID, CustomerName, CustomerAddress, CustomerPostalCode)

3.1.2. Mapping a **composite attribute** (Figure 4-9)

3.1.2. Composite attributes: Use only their simple component attributes

(a) CUSTOMER entity type with composite attribute



(b) CUSTOMER relation with address detail

CUSTOMER					
<u>CustomerID</u>	CustomerName	CustomerStreet	CustomerCity	CustomerState	CustomerPostalCode

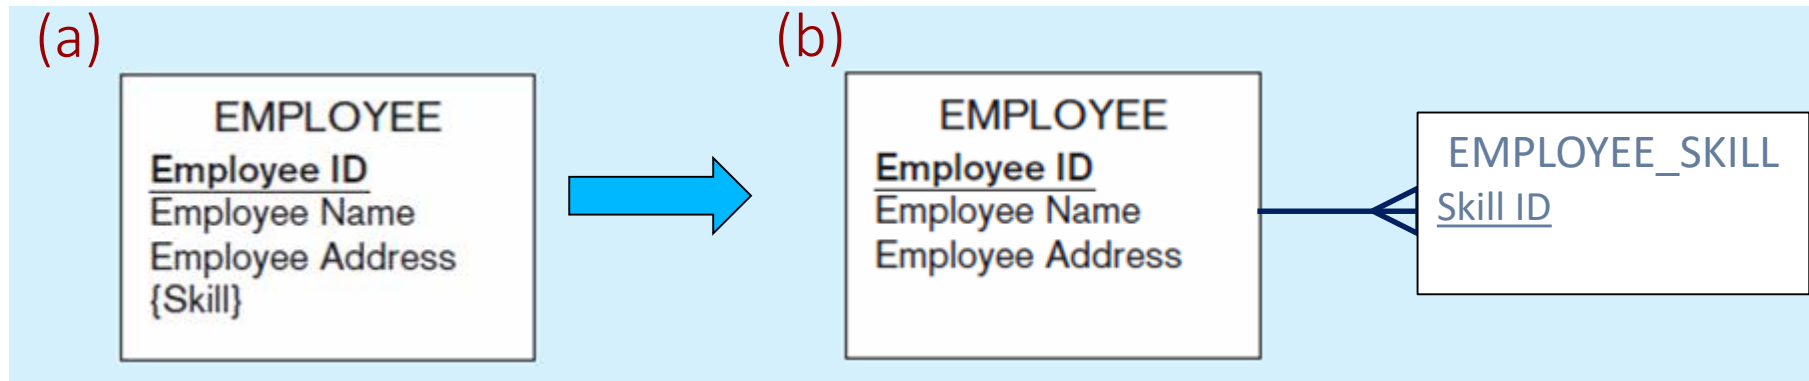
or

CUSTOMER (CustomerID, CustomerName, CustomerStreet, CustomerCity, CustomerState, CustomerPostalCode)

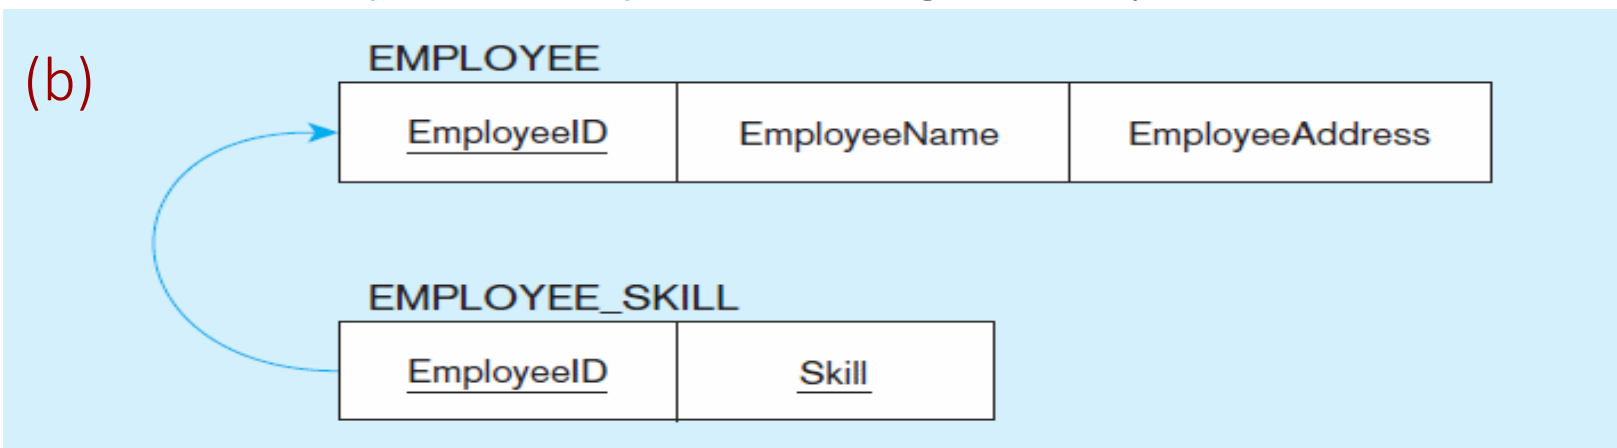
3.1.3. Mapping a multivalued attribute (Figure 4-10)

3.1.3. Multivalued Attribute:

Becomes a **separate relation** with a foreign key taken from the superior entity.




- Multivalued attribute becomes a **separate relation** with **foreign key**
- **One-to-many relationship** between original entity and new relation



Class Activity 4.2 (10 minutes)

2. Redesign the following entity, where every employee need to have at least one skill and every skill can be chosen by any employee. The information about skills like skill ID and name need to be stored in the database.

Explain why this need to be redesigned, and then convert your new ERD to the relations.

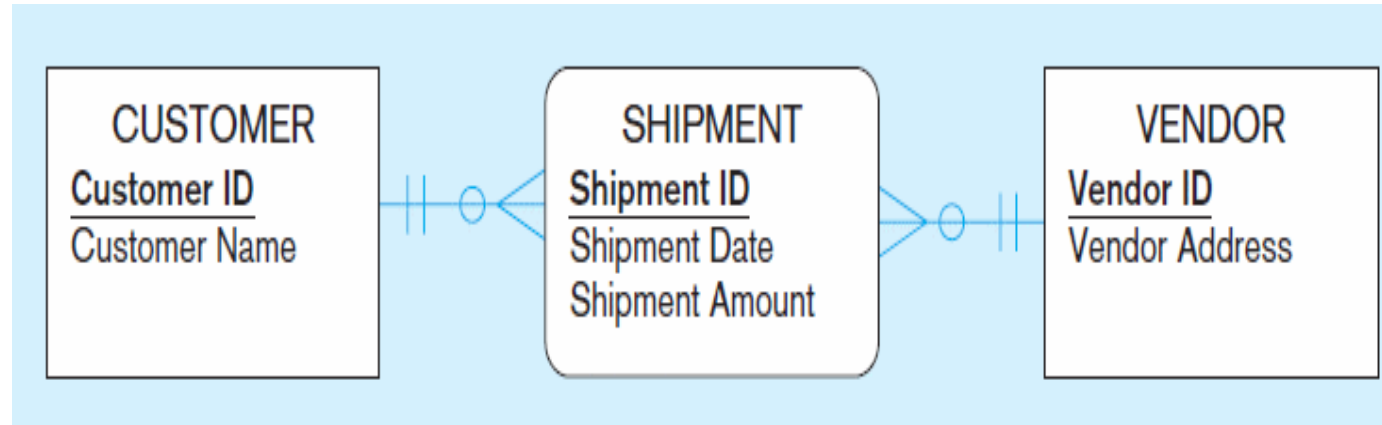


Solution to Class Activity 4.2:

EMPLOYEE
Employee ID
Employee Name
Employee Address
{Skill}

Class Activity 4.3 (5 minutes)

3. Convert the associative entity in this ERD to a relation?



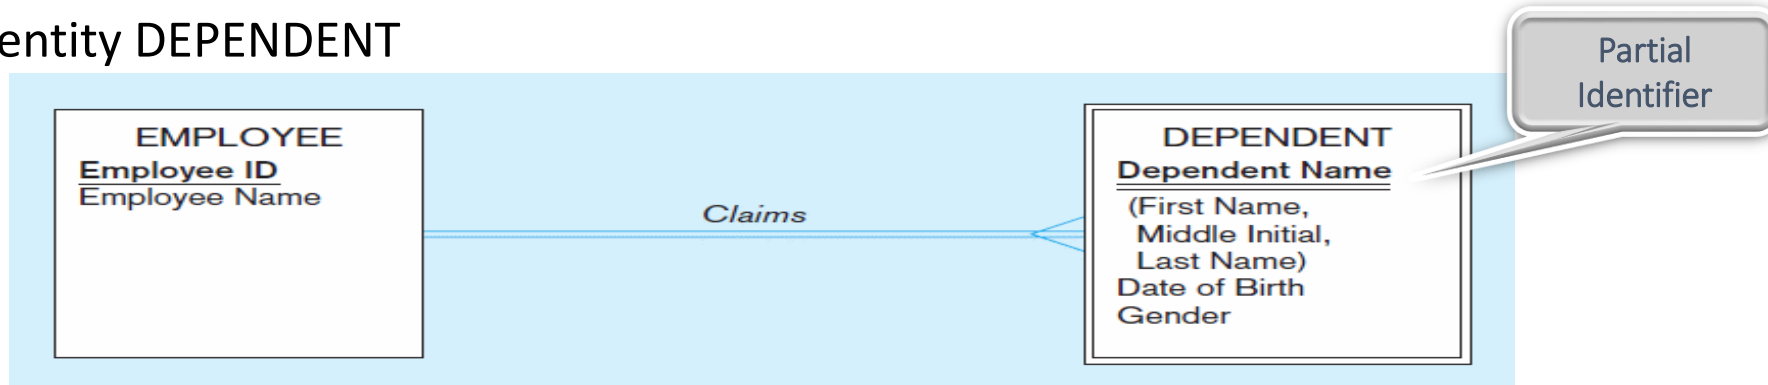
3.2. Mapping Weak Entities

- Becomes a **separate relation** with a **foreign key** taken from the superior entity
- **Primary key composed of:**
 - Partial identifier of weak entity
 - Primary key of identifying relation (strong entity)

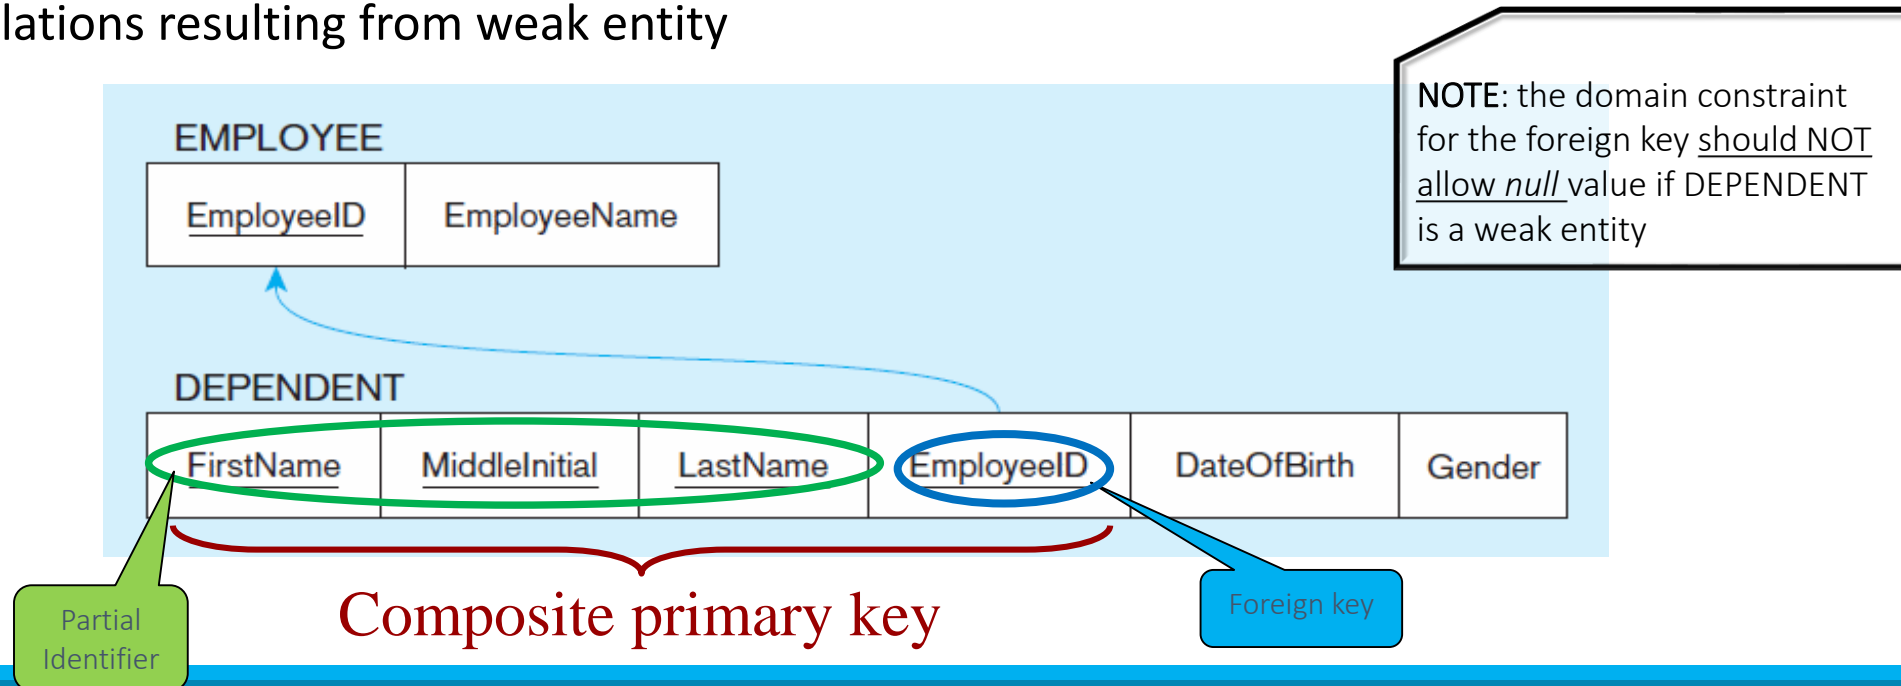
NOTE: Foreign keys can have null values, **but** the domain constraint for the foreign key should NOT allow null value if DEPENDENT is a **weak** entity or an **associative** entity, or is related to a **mandatory cardinality**.

Example of mapping a weak entity (Figure 4-11)

a) Weak entity DEPENDENT



b) Relations resulting from weak entity

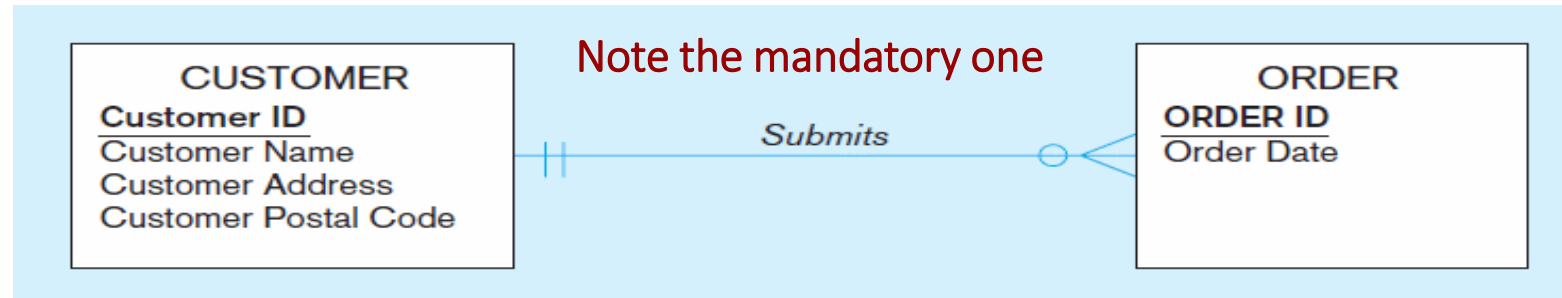


3.3. Mapping Binary Relationships

- 3.3.1. **One-to-Many**—Primary key on the **one** side becomes a foreign key on the **many** side
- 3.3.1. **Many-to-Many**—Create a ***new relation*** with the primary keys of the two entities as its primary key
- 3.3.1. **One-to-One**—Primary key on **mandatory** side becomes a foreign key on **optional** side

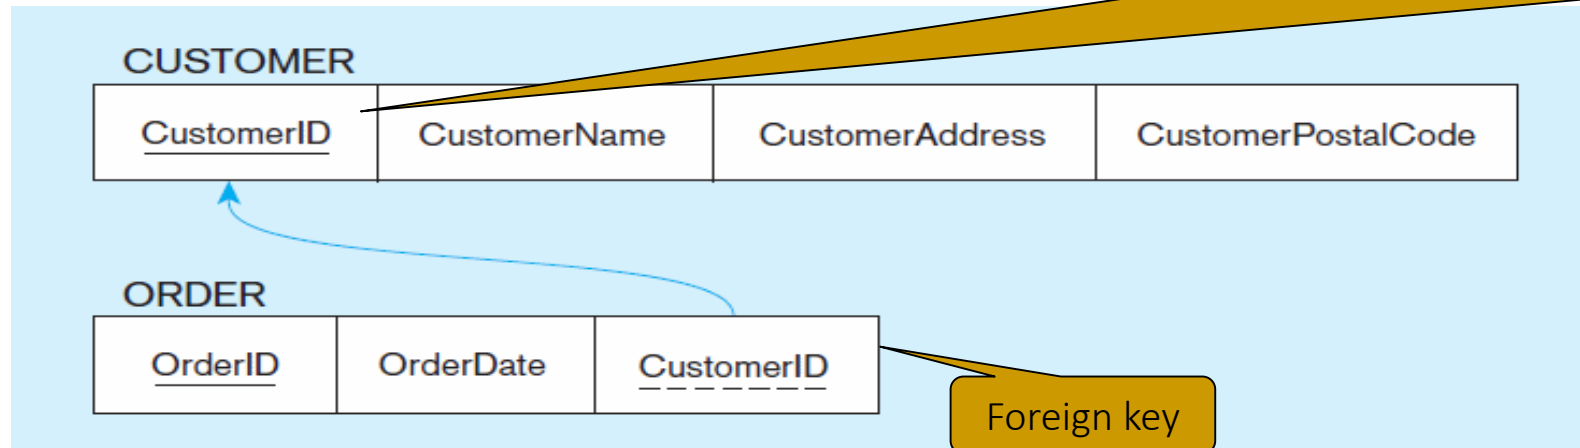
3.3.1. Mapping a 1:M relationship (Figure 4-12)

a) Relationship between customers and orders




Question: Can CustomerID in ORDER relation has null values?

b) Mapping the relationship

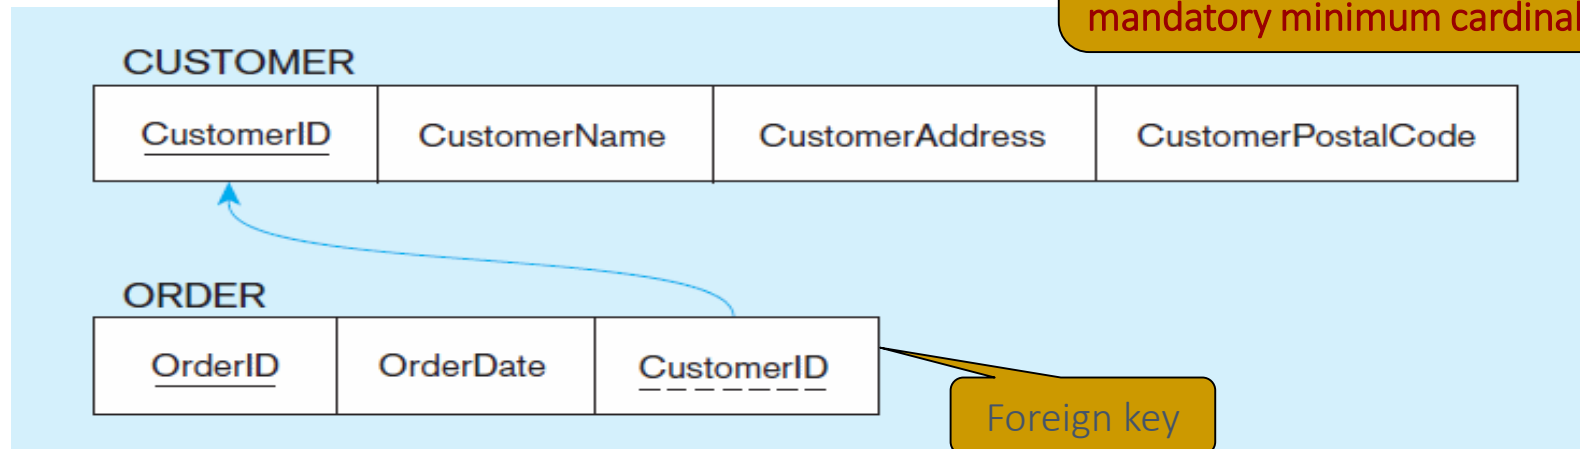


3.3.1. Mapping a 1:M relationship (Figure 4-12)

a) Relationship between customers and orders




b) Mapping the relationship



Again, no null value in the foreign key... this is because of the mandatory minimum cardinality.

3.3.1. Mapping a 1:M relationship- Other Format

a) Relationship between customers and orders



Again, no null value in the foreign key... this is because of the mandatory minimum cardinality.

b) Mapping the relationship

CUSTOMER (CustomerID, CustomerName, CustomerAddress, CustomerPostalCode)

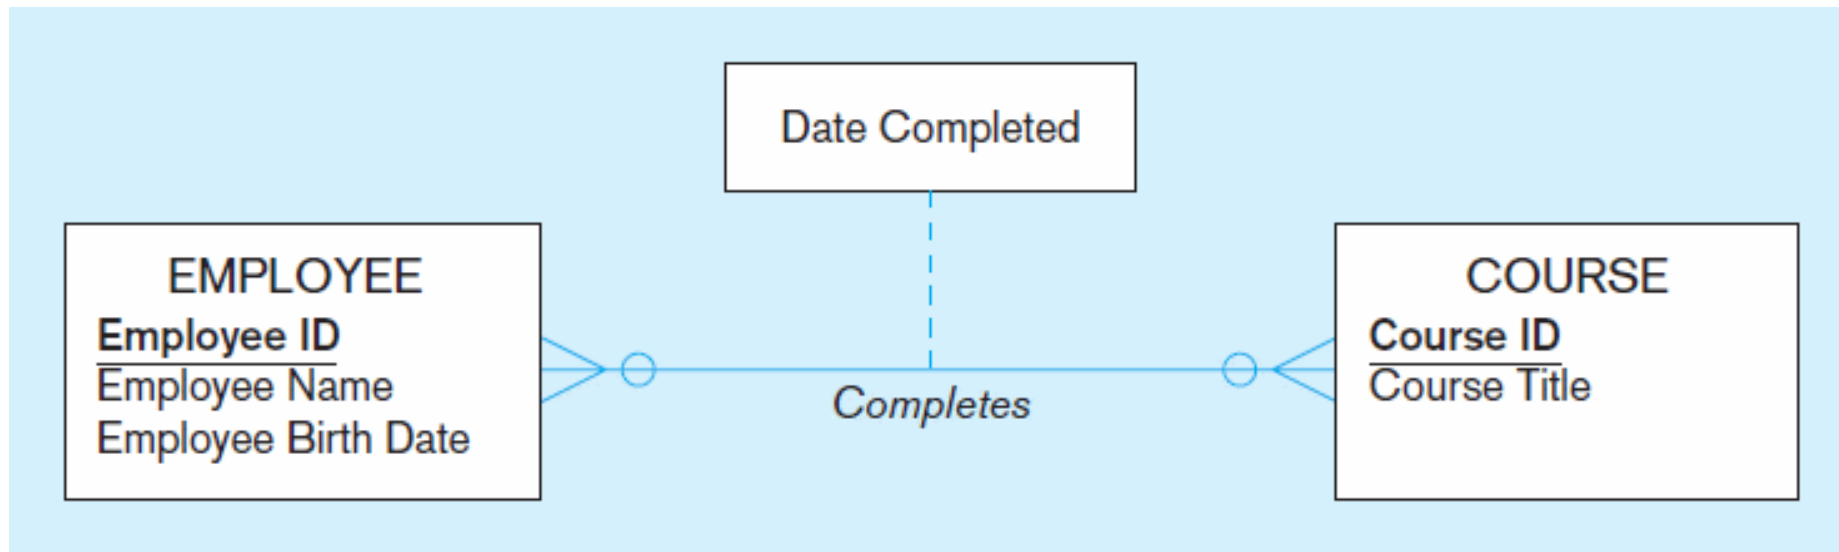
ORDER (OrderID, OrderDate, CustomerID*)

FK (CustomerID) references CUSTOMER

Foreign key

3.3.2. Mapping an M:N relationship (Figure 4-13)

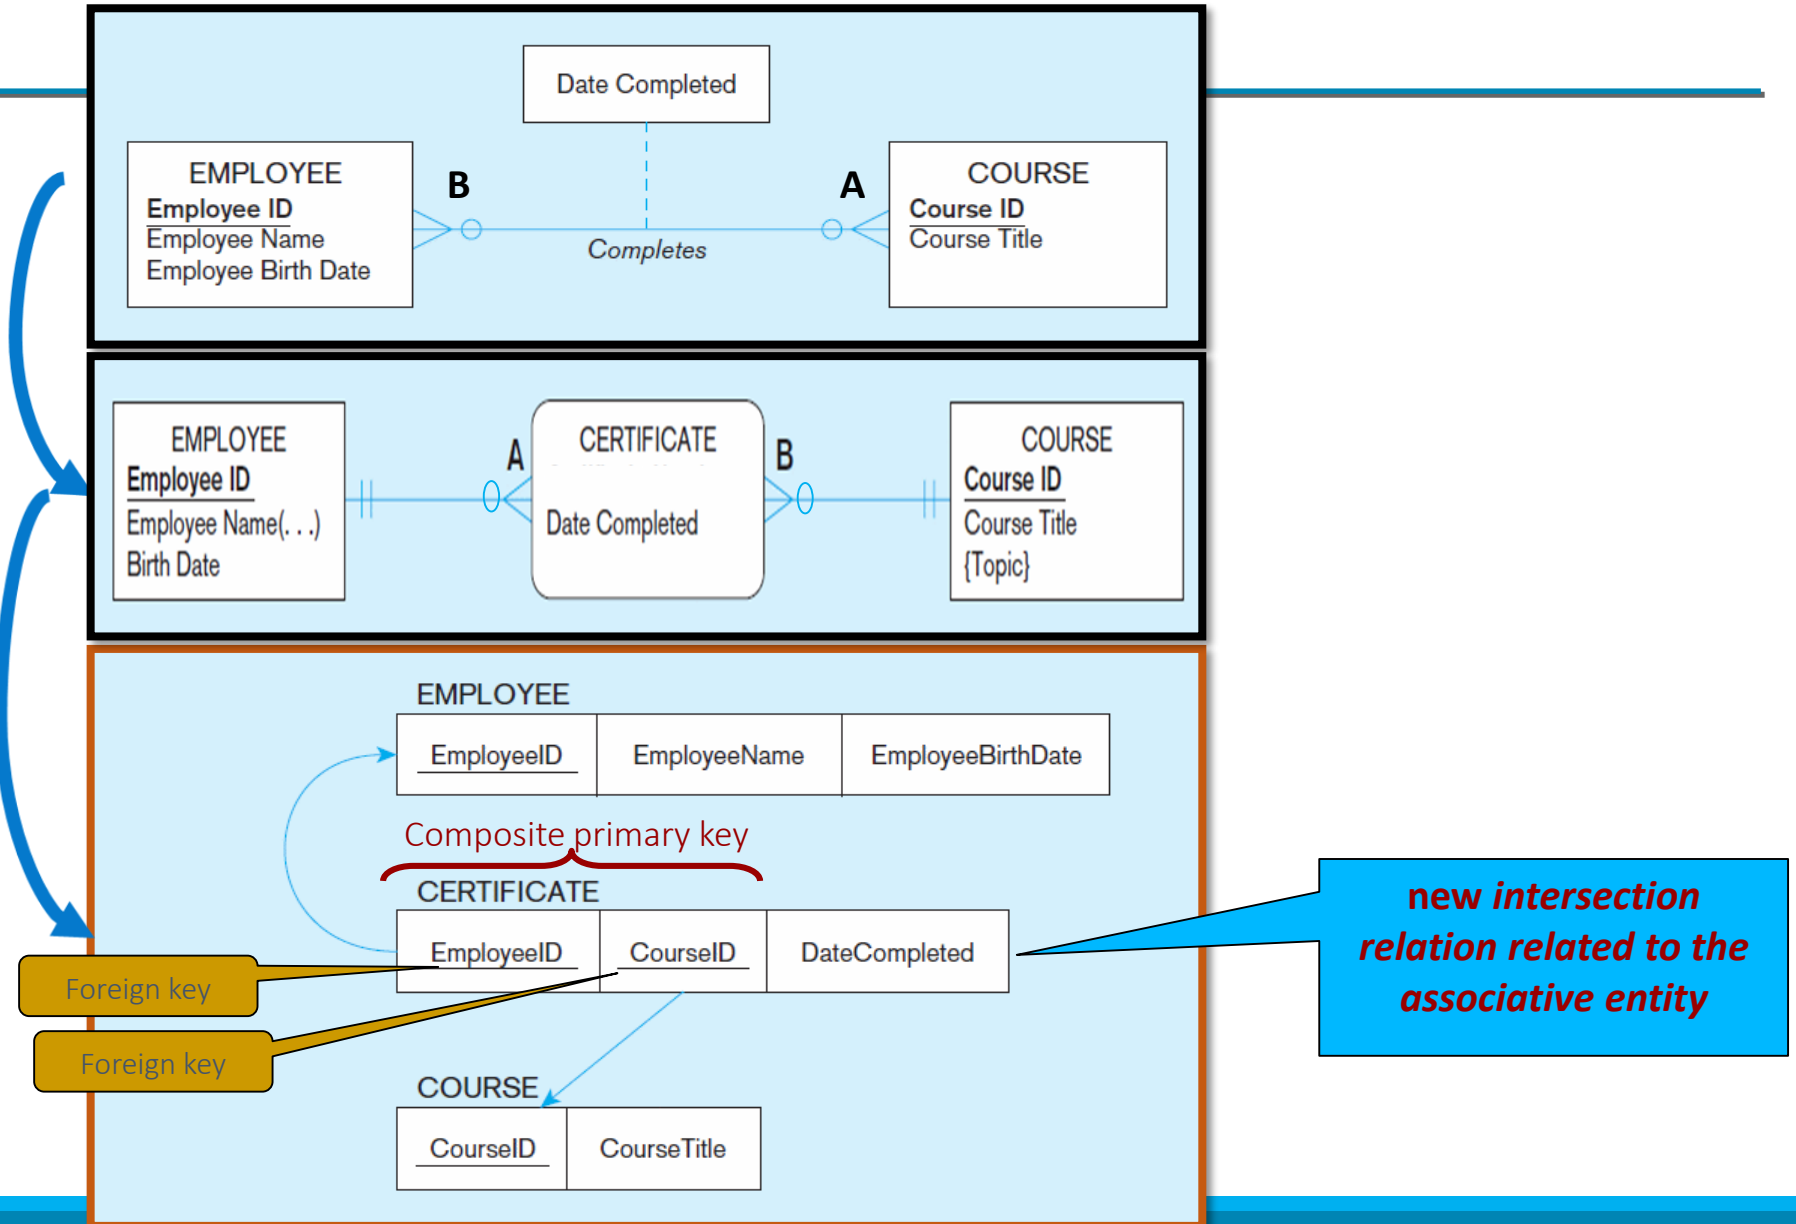
a) Completes relationship (M:N)



The *Completes* relationship will need to become a **separate relation**.

3.3.2. Mapping an M:N relationship (Figure 4-13) (cont.)

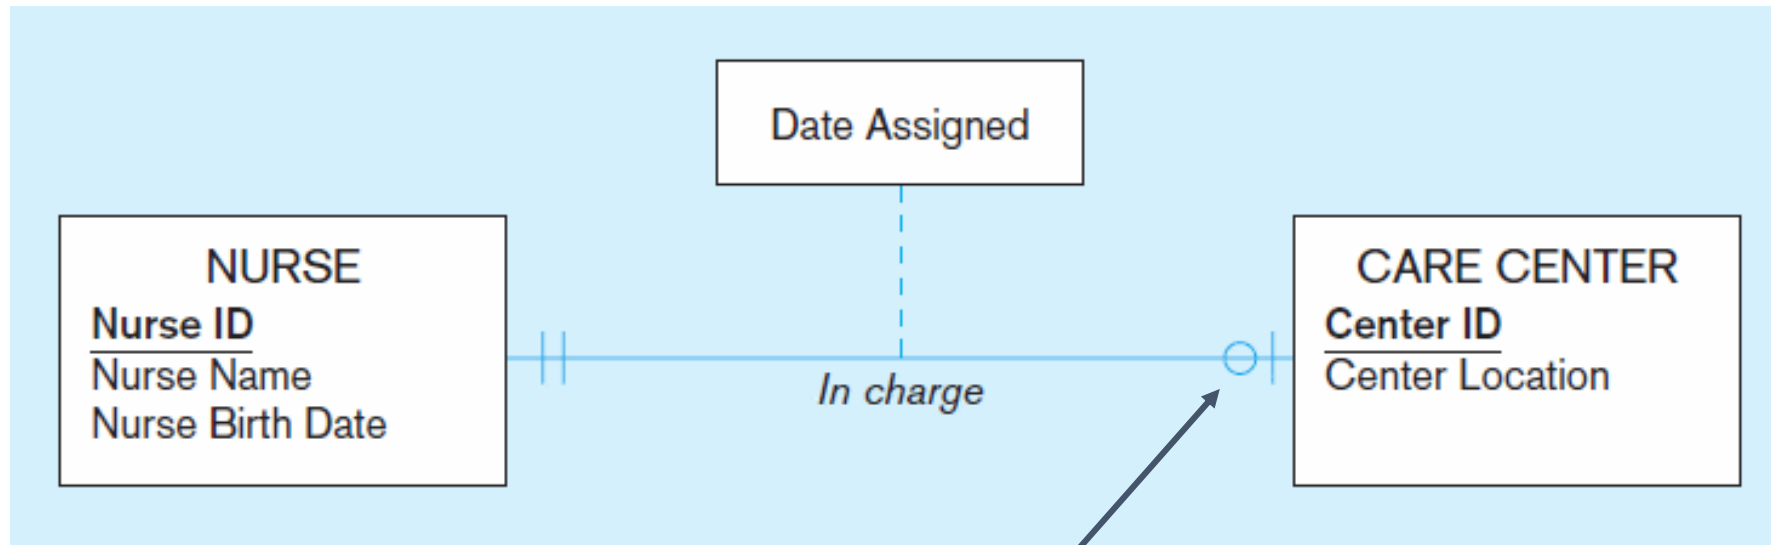
b) Three resulting relations



3.3.3. Mapping a binary 1:1 relationship (Figure 4-14)

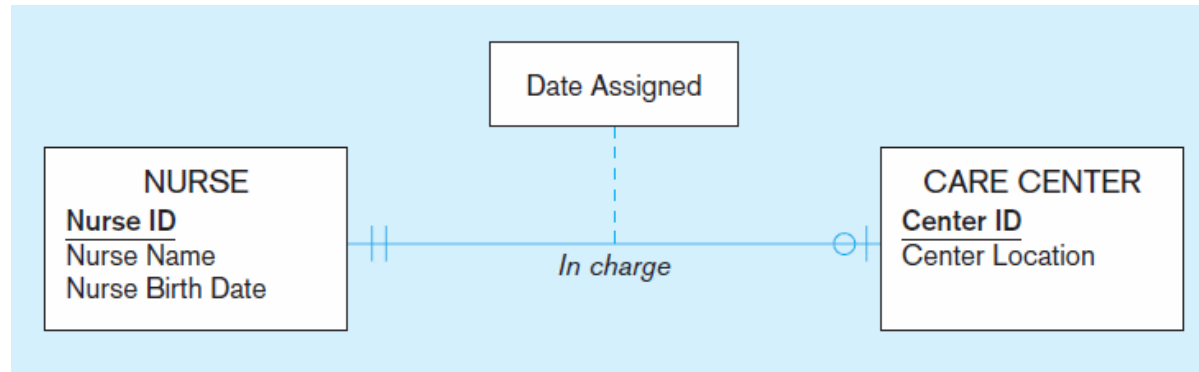
a) In charge relationship (1:1)

Rule: in 1:1 relationships, PK of the entity on the Mandatory side will be FK in the entity on the Optional side

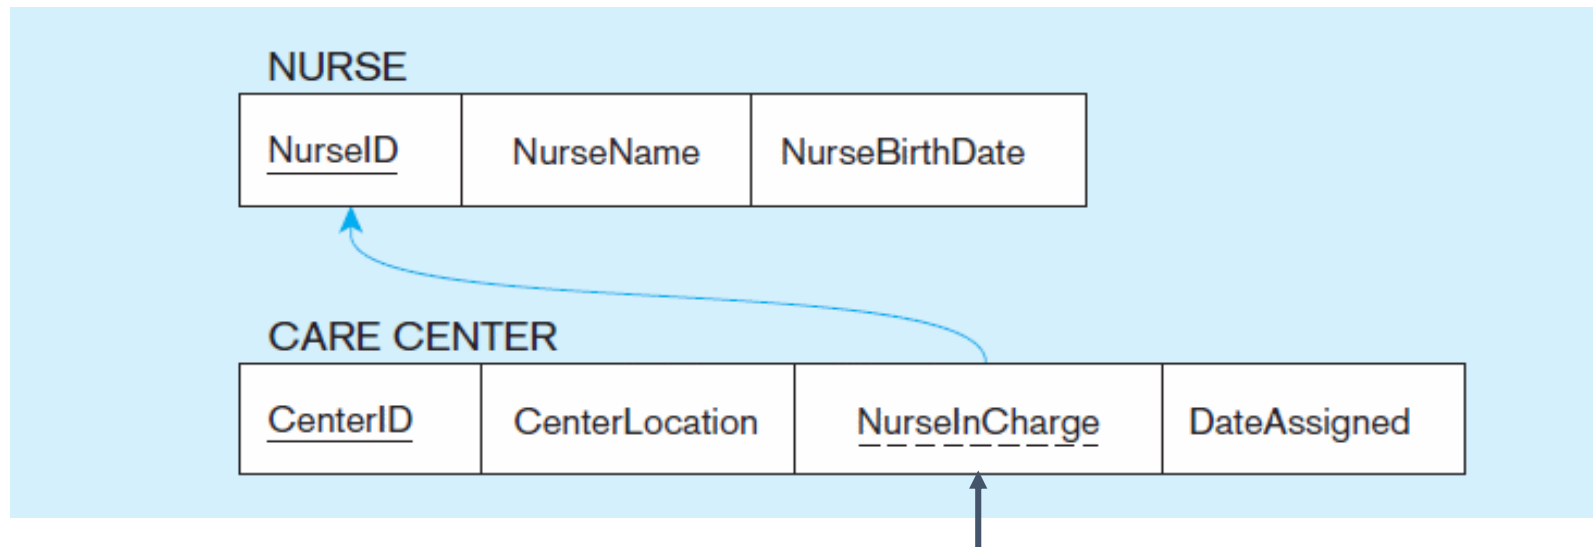


Often in 1:1 relationships, one direction is optional

3.3.3. Mapping a binary 1:1 relationship (Figure 4-14) (cont.)



b) Resulting relations



Foreign key goes in the relation on the **optional** side, matching the primary key on the mandatory side

3.4. Mapping Associative Entities

➤ Identifier Not Assigned

- Default primary key for the association relation is **composed** of the primary keys of the two entities
 - (as in M:N relationship)

➤ Identifier Assigned

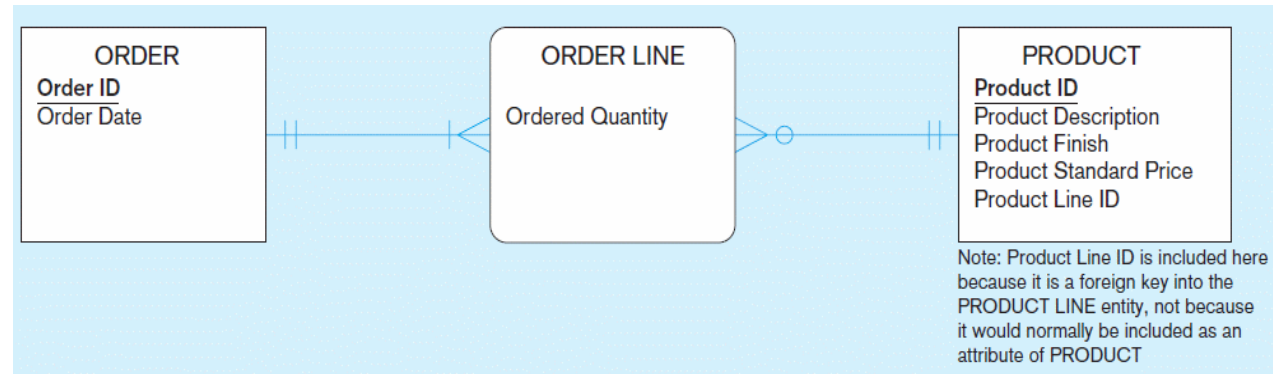
- It is natural and familiar to end-users
- Default identifier **may not be unique**

Example of mapping an associative entity (Figure 4-15)

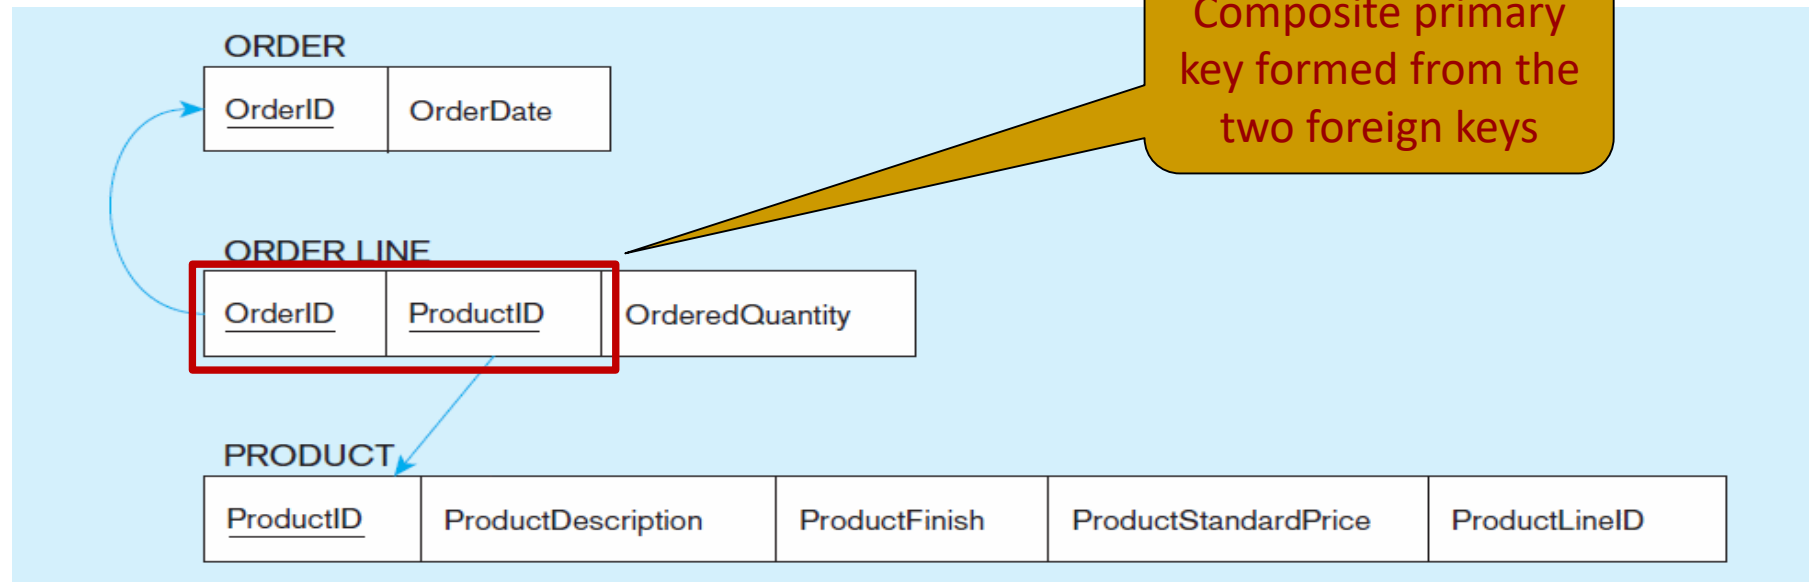
a) An associative entity



Example of mapping an associative entity (Figure 4-1) (cont.)

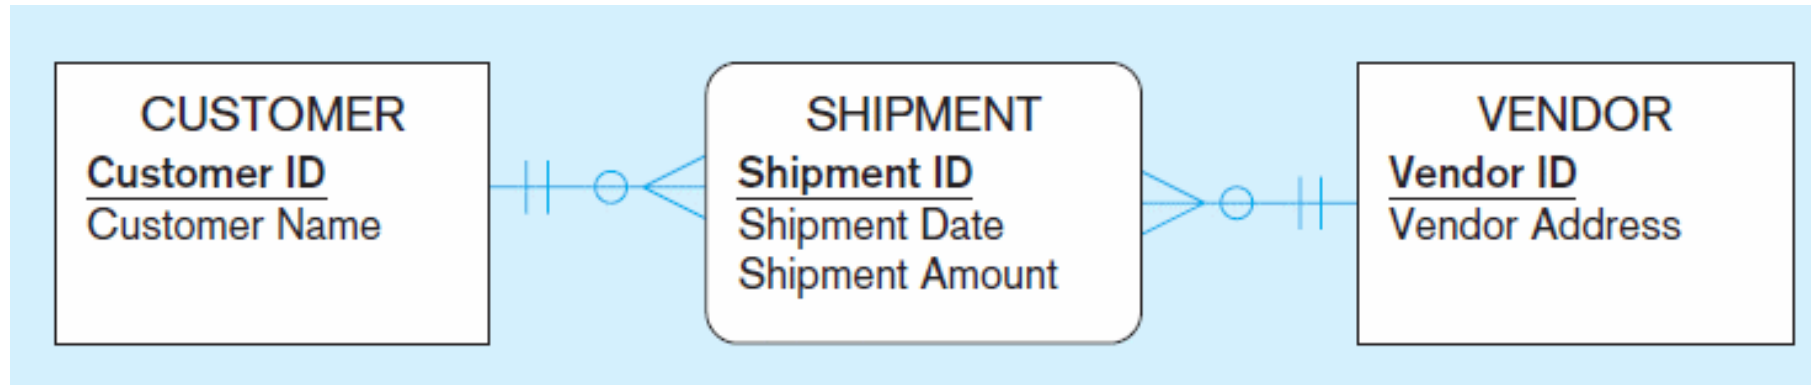


b) Three resulting relations

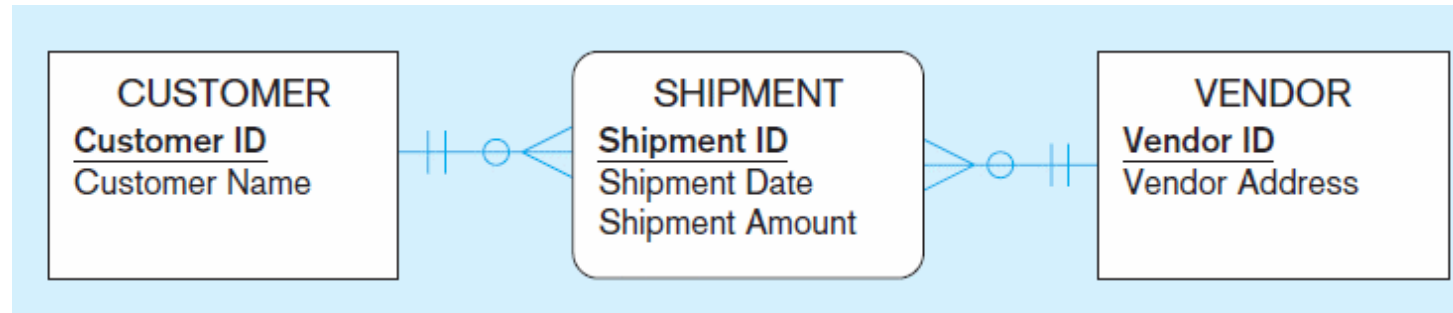


Example of mapping an associative entity with an identifier (Figure 4-16)

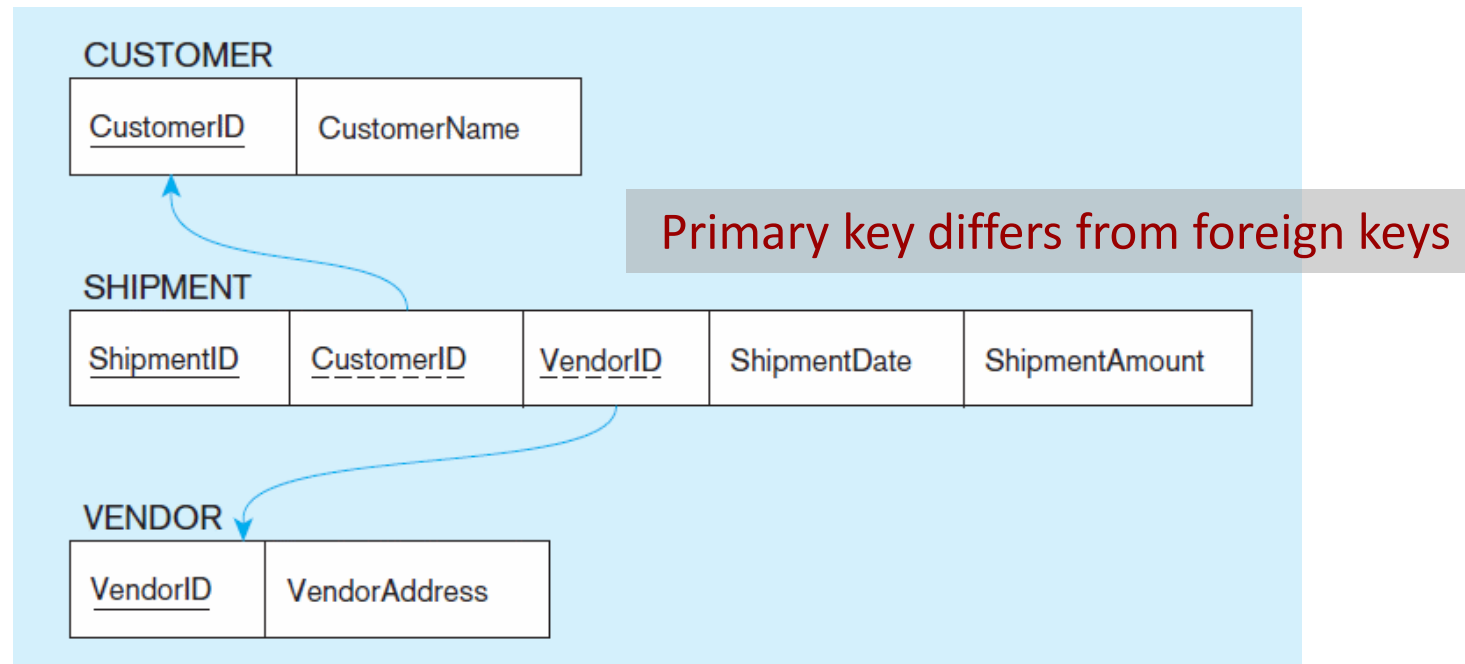
a) SHIPMENT associative entity



Example of mapping an associative entity with an identifier (Figure 4-16) (cont.)



b) Three resulting relations

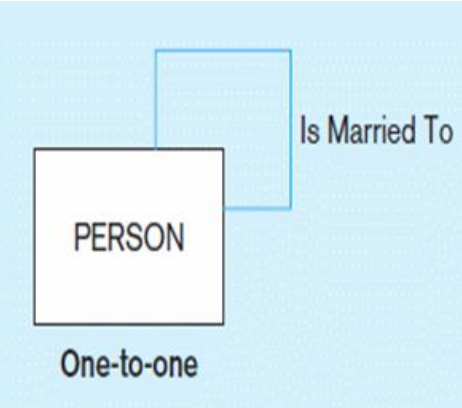


3.5. Mapping Unary Relationships

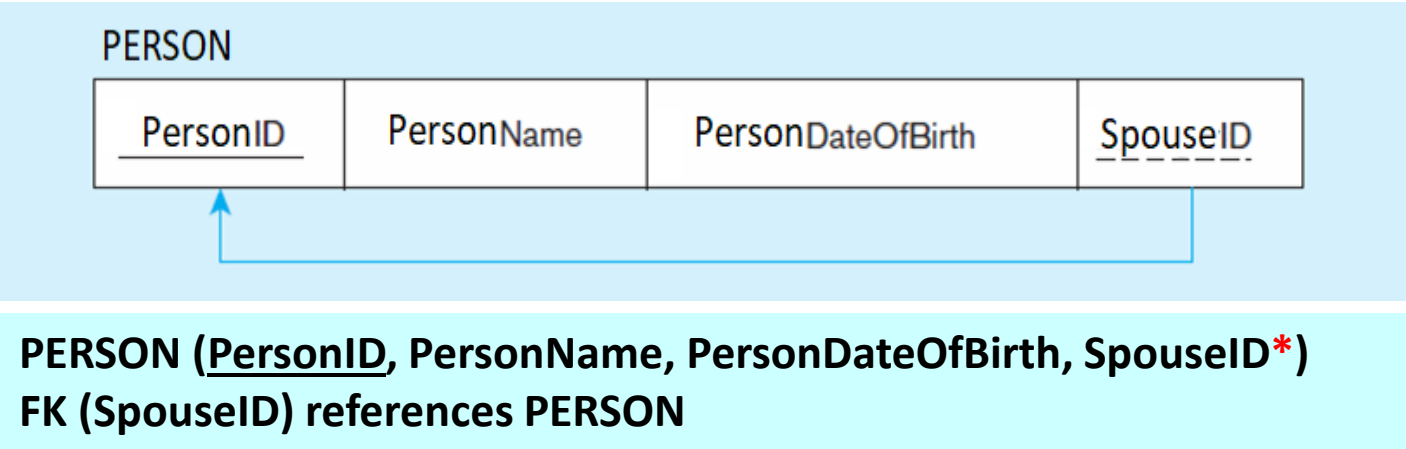
- **One-to-One and One-to-Many**—Recursive foreign key in the same relation
- **Many-to-Many**—Two relations:
 - One for the entity type
 - One for an associative relation in which the primary key has two attributes, both taken from the primary key of the entity

Mapping a unary 1:1 relationship

(a) **PERSON** entity with unary relationship



(b) **PERSON** relation with recursive foreign key

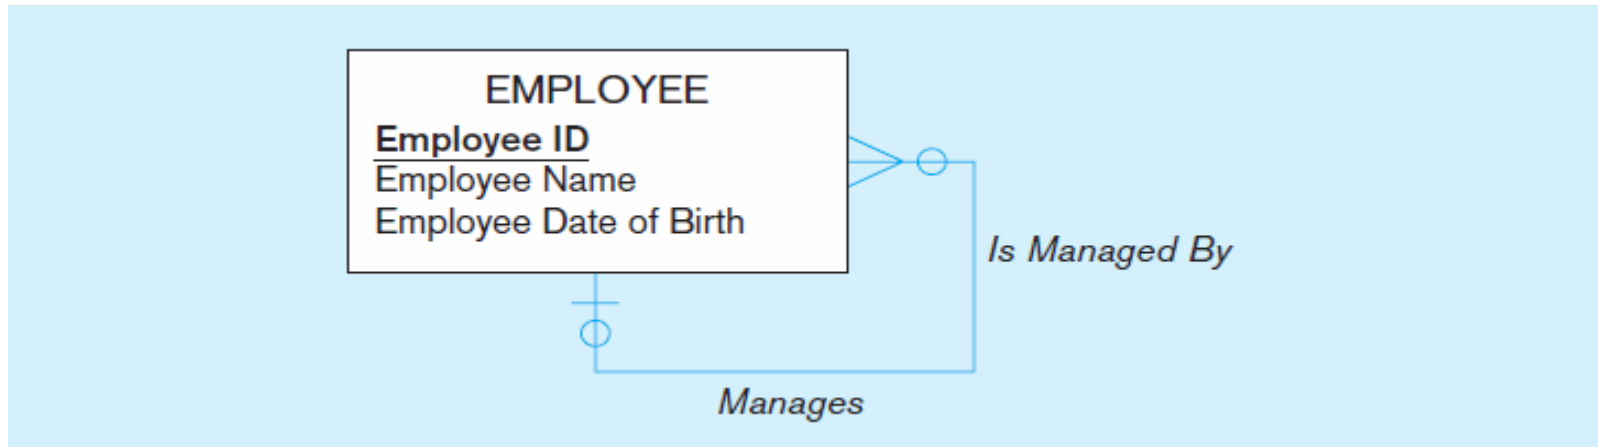


A table titled 'Employee' with four columns: 'PersonID', 'PersonName', 'PersonDateOfBirth', and 'Spouse ID'. A curved arrow labeled 'PK' points from the 'PersonID' column to the 'Spouse ID' column, which is labeled 'FK'. The data rows are as follows:

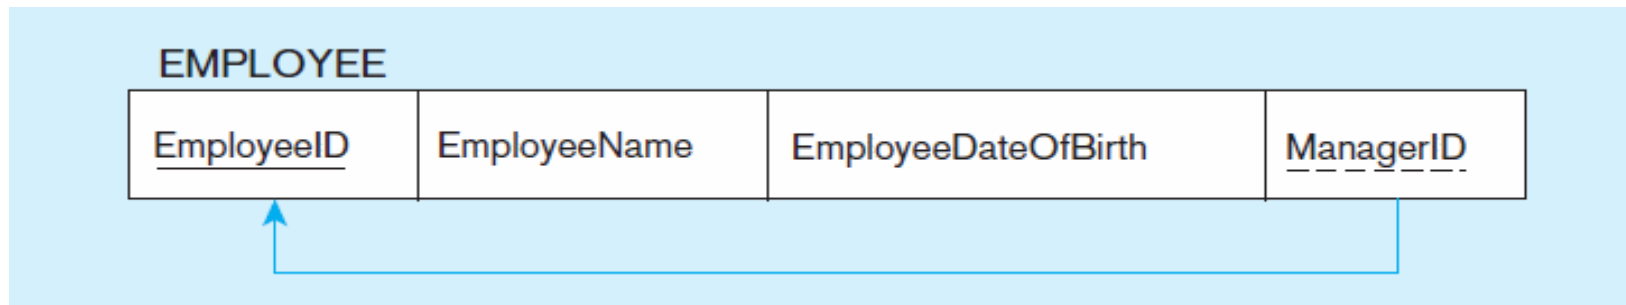
PersonID	PersonName	PersonDateOfBirth	Spouse ID
1234587	Jack	1/1/1980	2387468
3459087	Michael	5/2/1990	
2387468	Sara	5/8/1975	
5745321	Fahimeh	9/2/1983	
8743836	Ricky	5/11/1992	

Mapping a unary 1:N relationship (Figure 4-17)

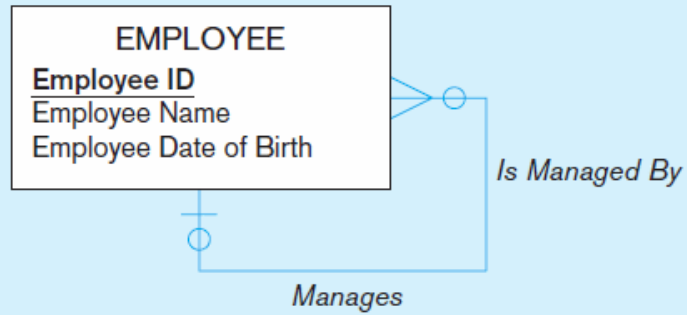
(a) **EMPLOYEE** entity with unary relationship



(b) **EMPLOYEE** relation with recursive foreign key



EMPLOYEE (EmployeeID, EmployeeName, EmployeeDateOfBirth, ManagerID*)
FK (ManagerID) references EMPLOYEE



EMPLOYEE

<u>EmployeeID</u>	EmployeeName	EmployeeDateOfBirth	<u>ManagerID</u>
-------------------	--------------	---------------------	------------------

EMPLOYEE (EmployeeID, EmployeeName, EmployeeDateOfBirth, ManagerID*)
FK (ManagerID) references EMPLOYEE

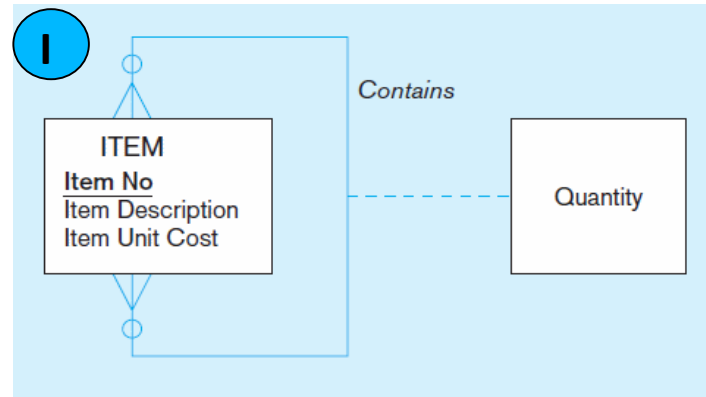
PK

FK

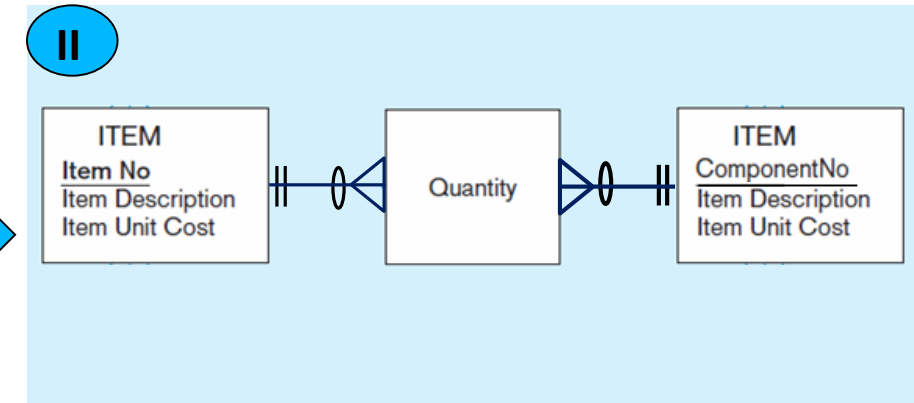
Employee_ID	Employee_Name	Employee_DateOfBirth	Manager_ID
1123	Sara	1.1.2000	7892
1456	Jake	1.1.2000	7892
7892	Fahimeh	1.1.1970	1245
1245	Julia	1.1.1980	...
...

Mapping a unary M:N relationship (Figure 4-18)

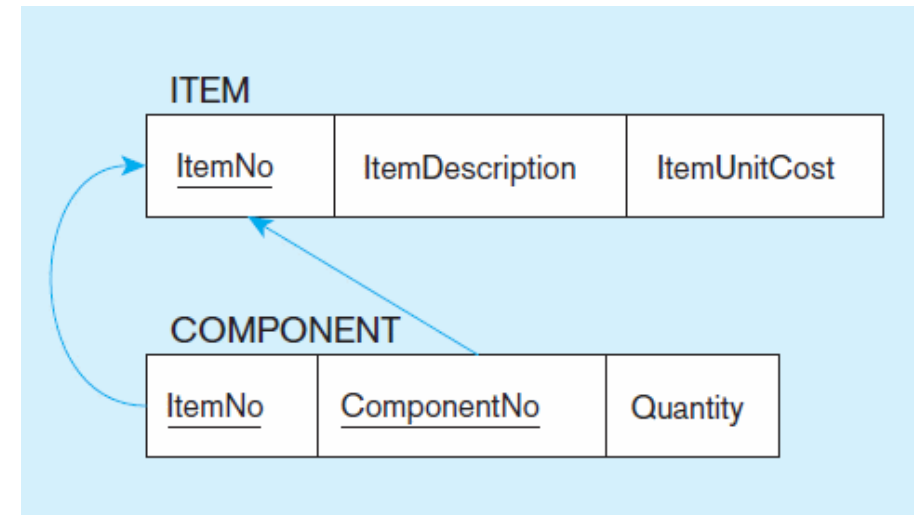
➤ Bill-of-materials relationships (M:N)



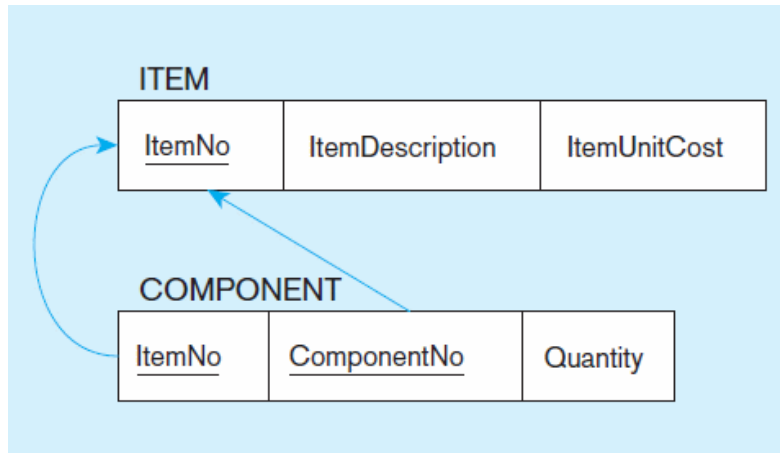
ERD II is just provided to clarify ERD I.
ERD II is not a standard ERD



➤ ITEM and COMPONENT relations



Example



Item_No	Item_Description	Item_Unit_Cost
12	Wheel	50
13	Spoke	0.5
14	Rim	30
15	Valve	5

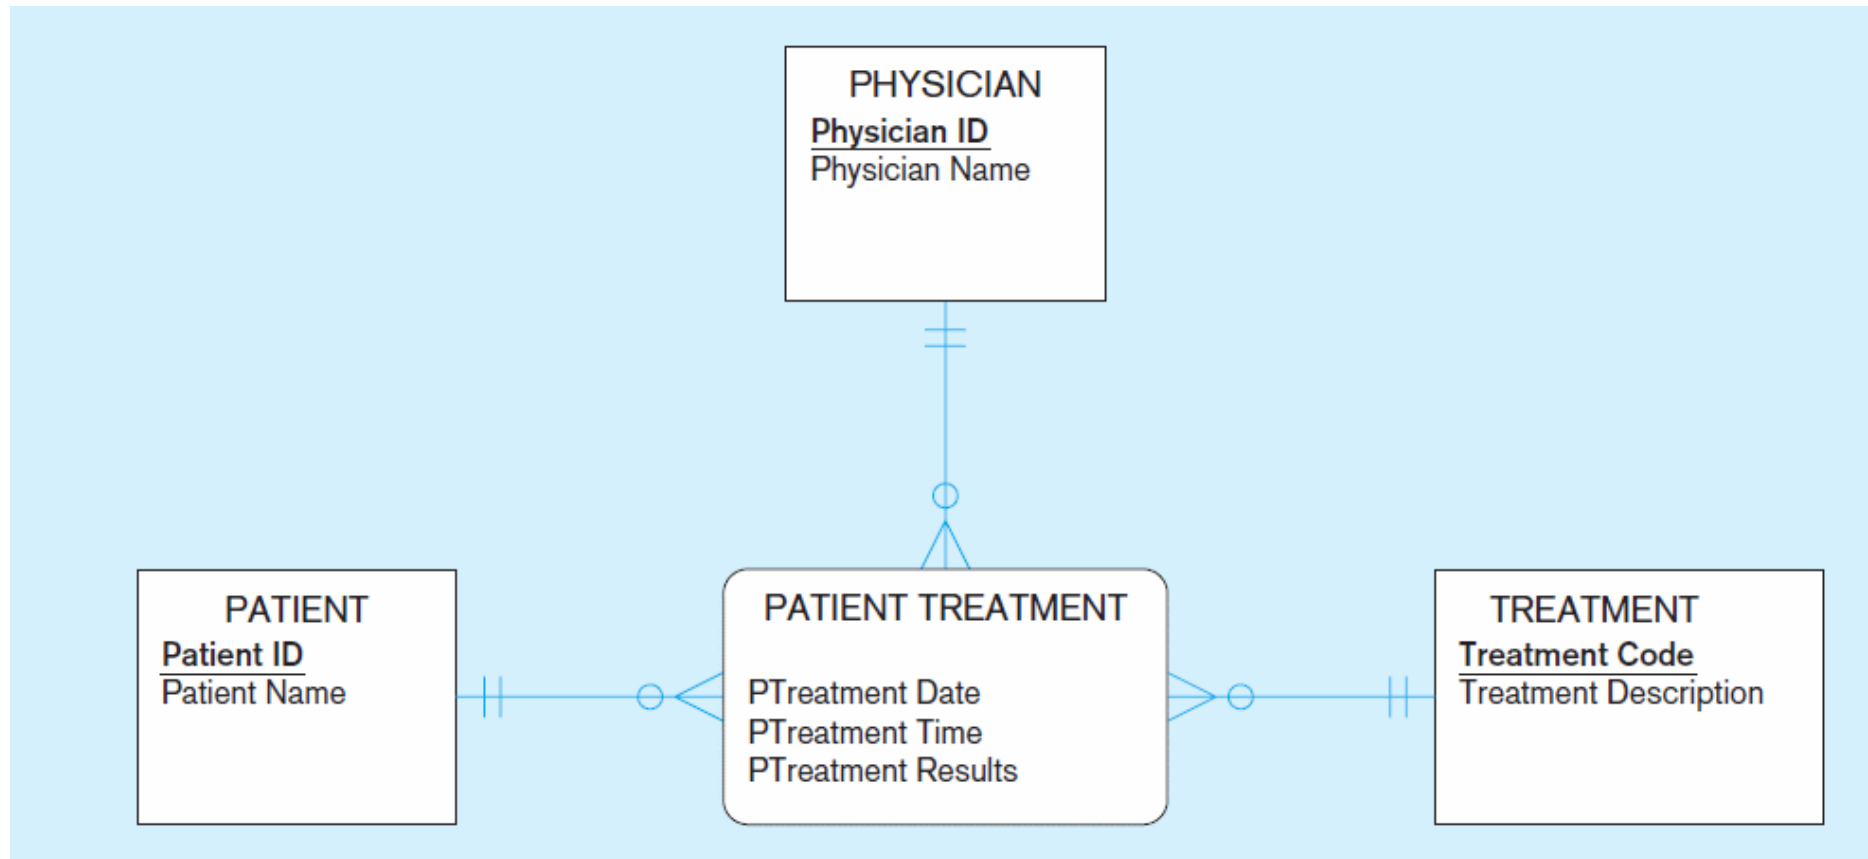
Item_No	Component_No	Quantity
12	13	30
12	14	1
12	15	1

3.6. Mapping Ternary (and n-ary) Relationships

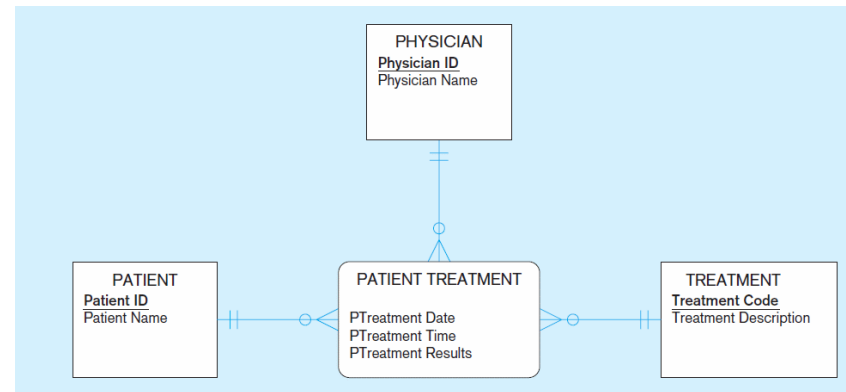
- **One relation for each entity and one for the associative entity**
- **Associative entity has foreign keys to each entity in the relationship**

Mapping a ternary relationship (Figure 4-19)

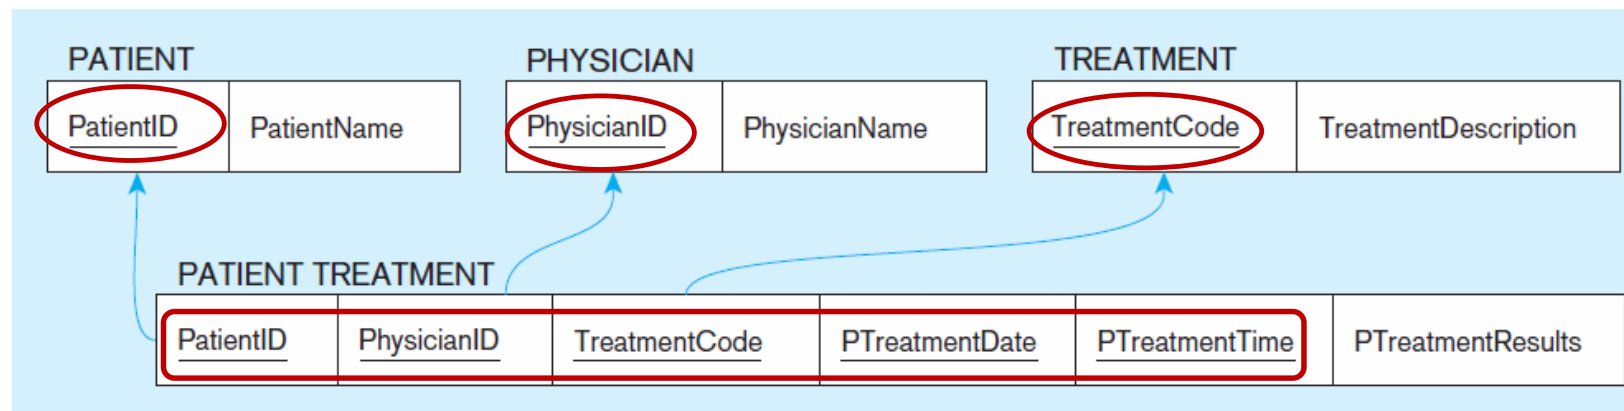
a) PATIENT TREATMENT Ternary relationship with associative entity



Mapping a ternary relationship (Figure 4-19) (cont.)



b) Mapping the ternary relationship PATIENT TREATMENT



Remember that the primary key **MUST** be unique.

This is why treatment **date** and **time** are included in the composite primary key.

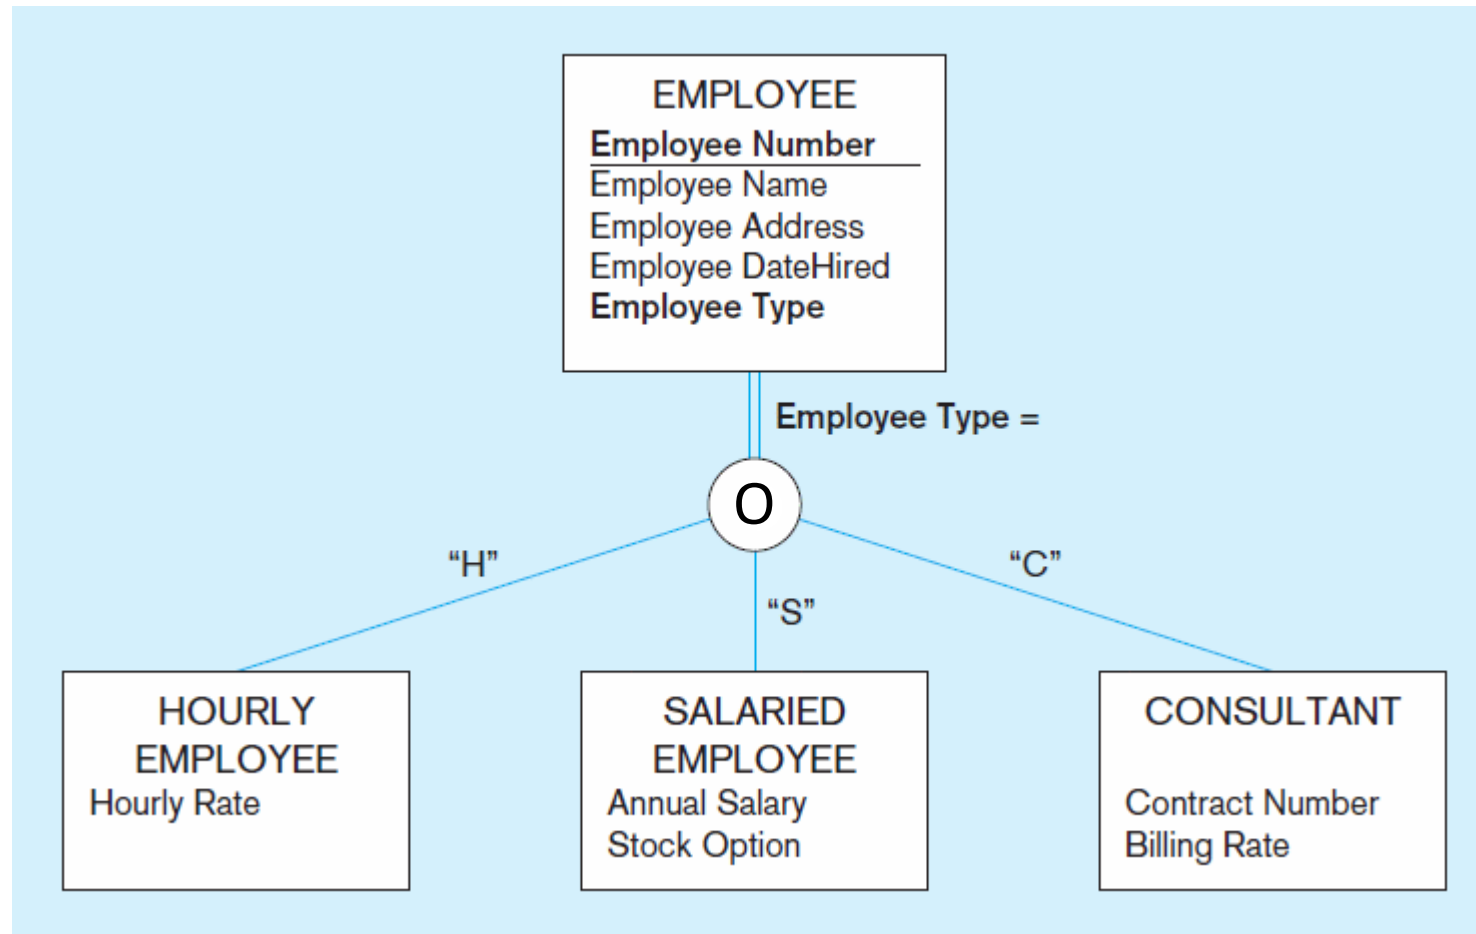
But this makes a very cumbersome key...

It would be better to create a **surrogate** key like Treatment#.

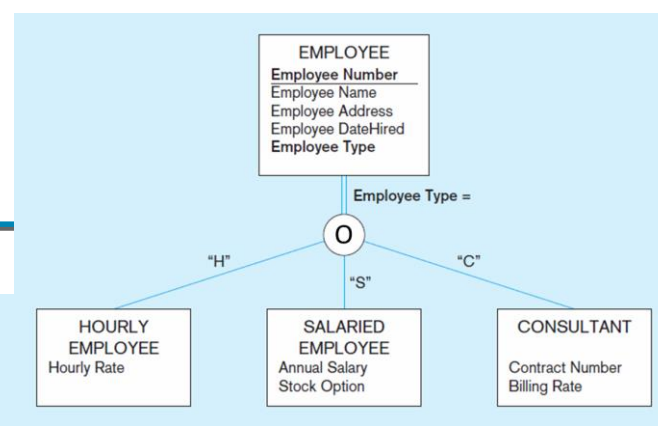
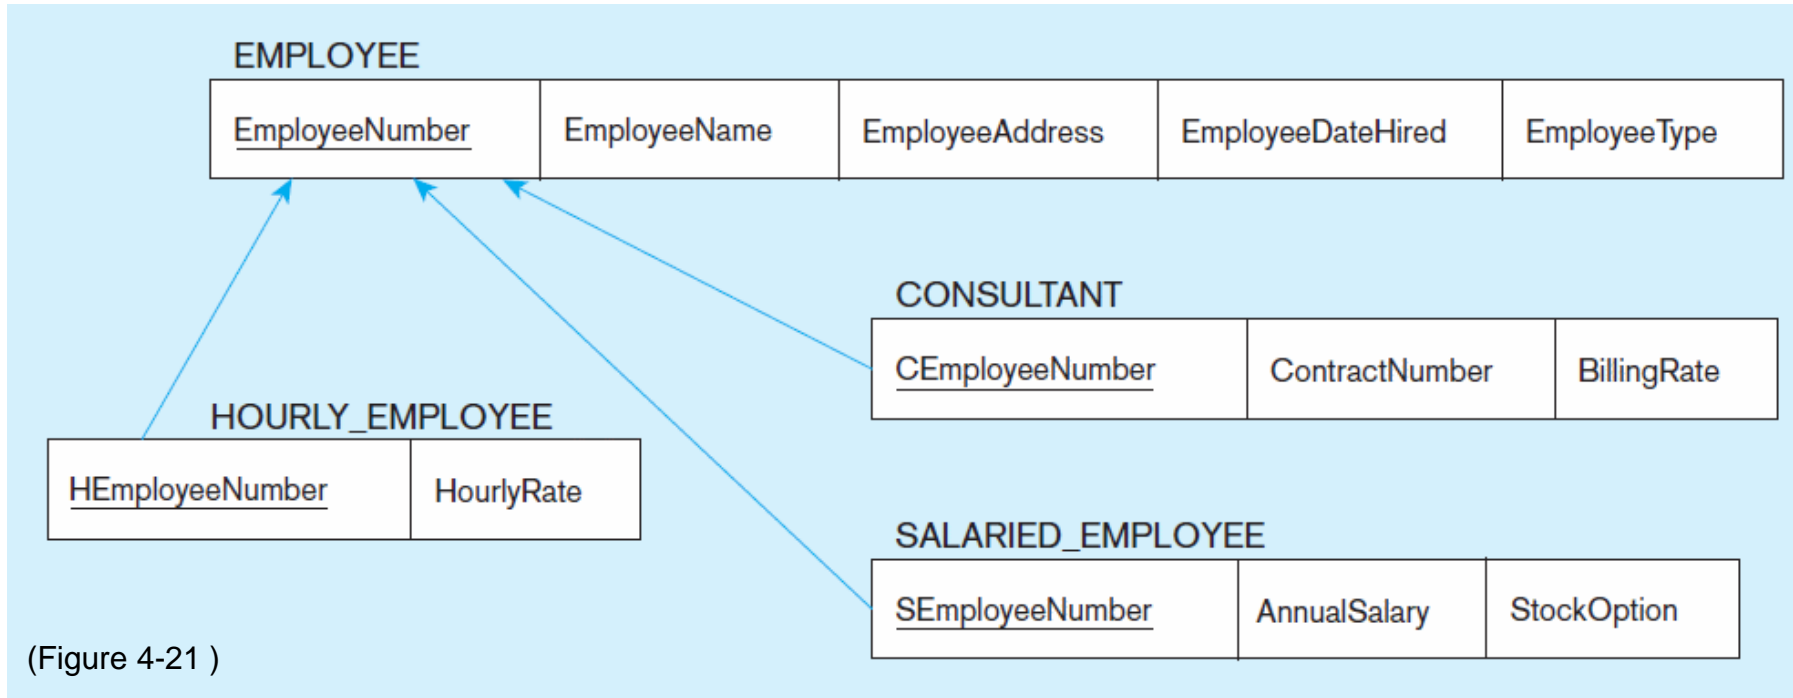
3.7. Mapping Supertype/Subtype Relationships

- One relation for supertype and for each subtype
- Supertype attributes (including **identifier** and subtype **discriminator**) go into supertype relation
- Subtype attributes go into each subtype; **primary key of supertype** relation also becomes **primary key of subtype relation**
- 1:1 relationship established between supertype and each subtype, with supertype as primary table

Supertype/subtype relationships (Figure 4-20)



Mapping supertype/subtype relationships to relations



These are implemented as one-to-one relationships.

Note: This is the best method to map supertype/subtypes to relations. There are other two methods that will be discussed in the related tutorial.

Introducing a subtype discriminator (overlap rule)

Employee_Number	Employee_Name	Address	Date_Hire	Employee_type (H? S? C?)
1123	Sara	UTS	1/1/2014	YNY
1456	Jake	32/50 ...	5/8/2013	NYN
7892	Fahimeh	12/97 ...	2/3/2013	NNY

Composite
Attribute



HEmployeeNumber	HourlyRate
1123	80

SEmployeeNumber	AnnualSalary	StockOption
1456	70000	0.2

CEmployeeNumber	ContractNumber	BillingRate
7892	9856	50
1123	9812	30

Summary

- List properties of relations
- Transform E-R and EER diagrams to relations
- Create tables with entity and relational integrity constraints

Next Lecture...

1. Terms to know to Do Normalization

- 1.1. Functional Dependencies
- 1.2. Keys: Super-key, Candidate key and Primary Key
- 1.3. Determining Candidate Keys from FDs
- 1.4. Partial Functional Dependencies
- 1.5. Transitive Functional Dependencies

2. Data Normalization and Well-Structured Relations

3. Steps in normalization

4. First Normal Form

5. Second Normal Form

6. Third Normal Form



This work is protected by United States copyright laws and is provided solely for the use of instructors in teaching their courses and assessing student learning. Dissemination or sale of any part of this work (including on the World Wide Web) will destroy the integrity of the work and is not permitted. The work and materials from it should never be made available to students except by instructors using the accompanying text in their classes. All recipients of this work are expected to abide by these restrictions and to honor the intended pedagogical purposes and the needs of other instructors who rely on these materials.